

# Generative Adversarial Networks – Part II

Du Ang

April 28, 2018



VISION@OUC



# Recap

- Supervised Learning vs Unsupervised Learning
- Introduction to generative models
- Popular generative models in detail
  - Variational Autoencoders (VAE)
  - Generative Adversarial Networks (GAN)



## Recap:

# Supervised Learning vs Unsupervised Learning

## Supervised Learning

**Data:**  $(x, y)$

$x$  is data,  $y$  is label

**Goal:**

Learn a function to map  $x \rightarrow y$

**Examples:** Regression,  
classification, object detection,  
semantic segmentation, image  
captioning, etc.

## Unsupervised Learning

**Data:**  $x$

Just data, no labels!

**Goal:**

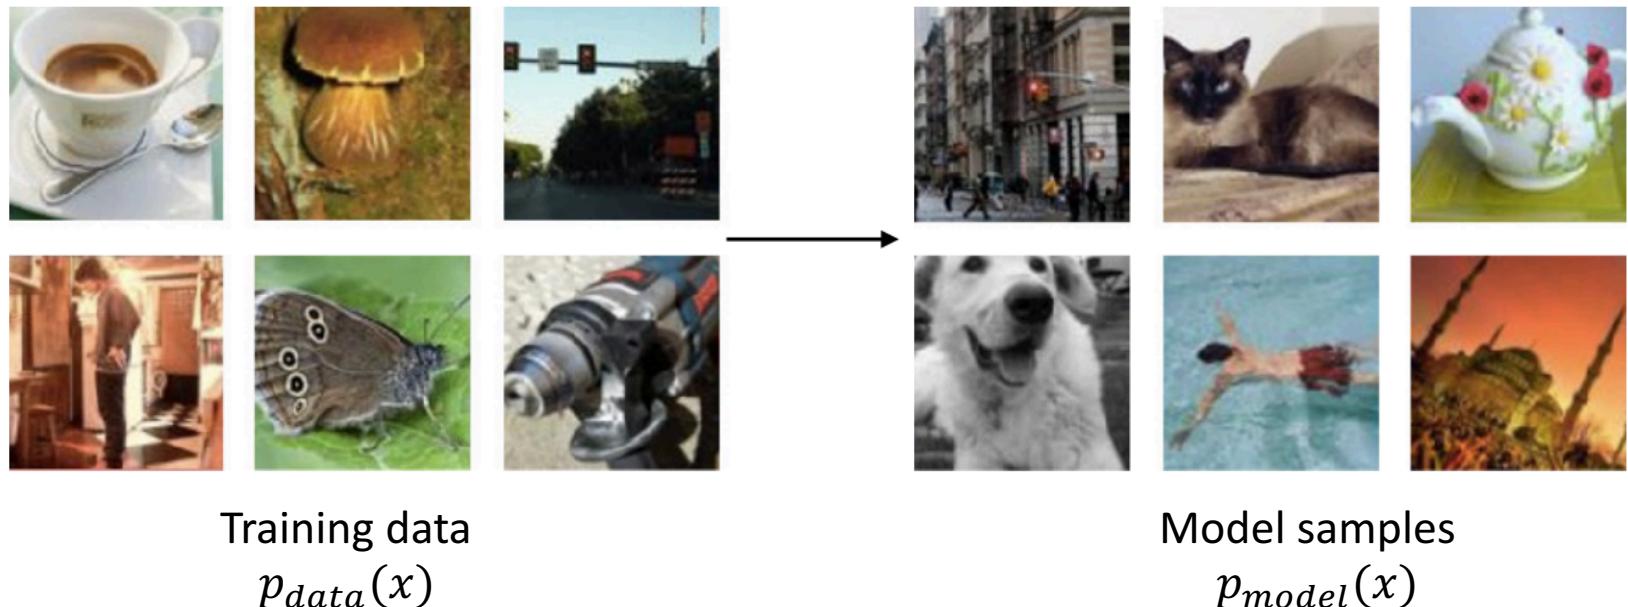
Learn some underlying  
hidden *structure* of the data

**Examples:** Clustering,  
dimension reduction,  
feature learning,  
density estimation, etc.



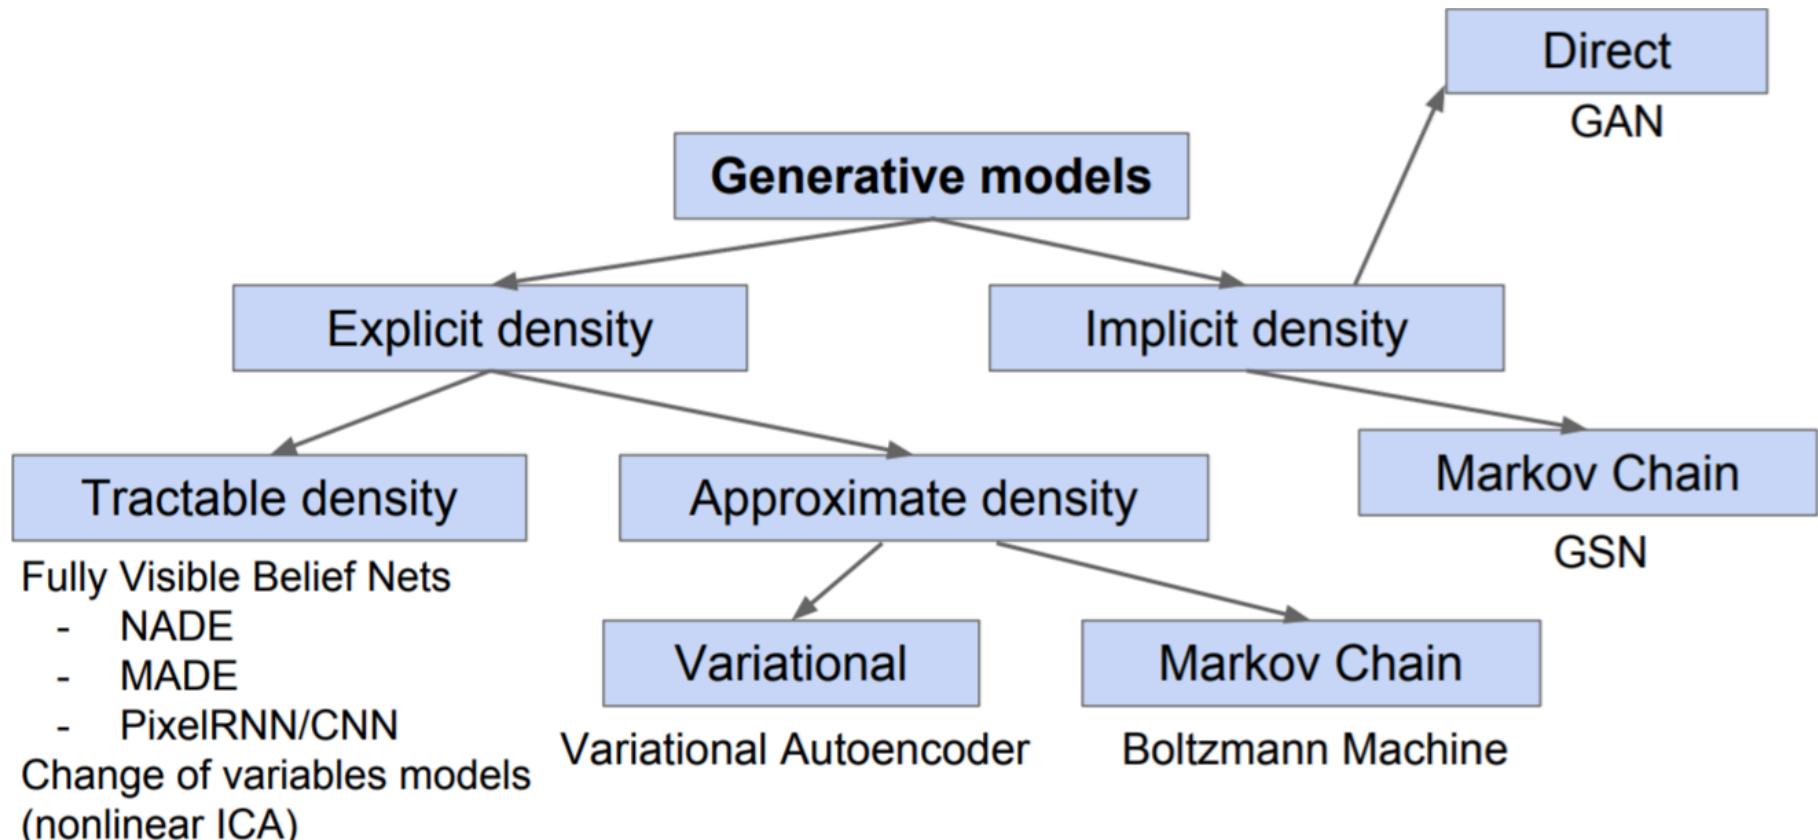
# Recap: Generative Models

Takes a training set, consisting of samples drawn from a distribution  $p_{data}$ , and learns to represent an estimate of that distribution with  $p_{model}$ .





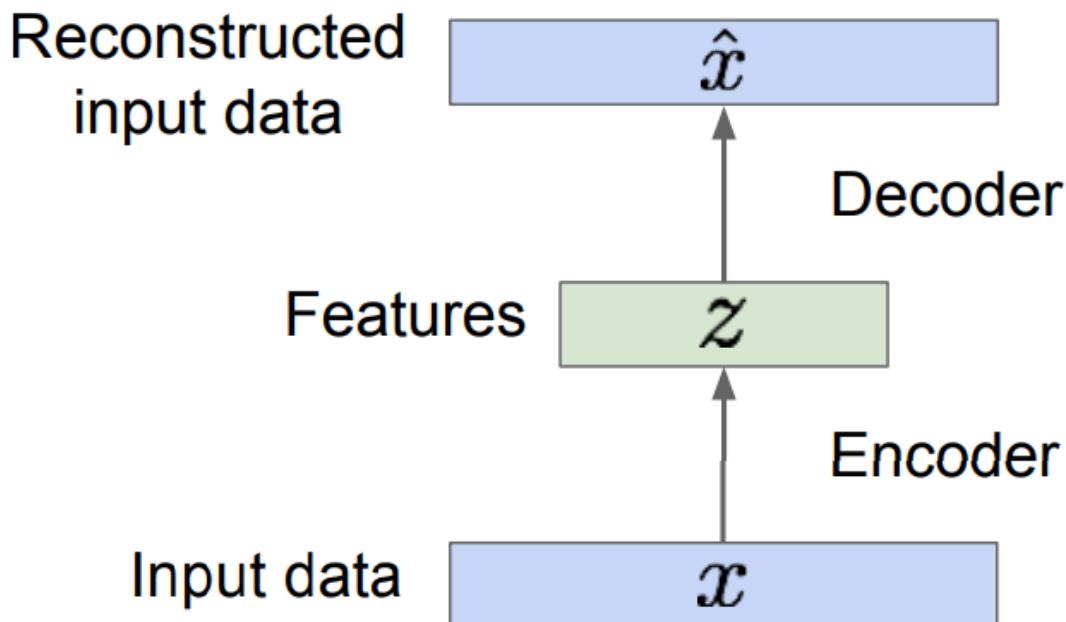
# Recap: Taxonomy of Generative Models





# Recap: Autoencoders

- Autoencoders can reconstruct data, and can learn features to initialize a supervised model.
- Features capture factors of variation in training data.





# Recap: Variational Autoencoders

## Data likelihood

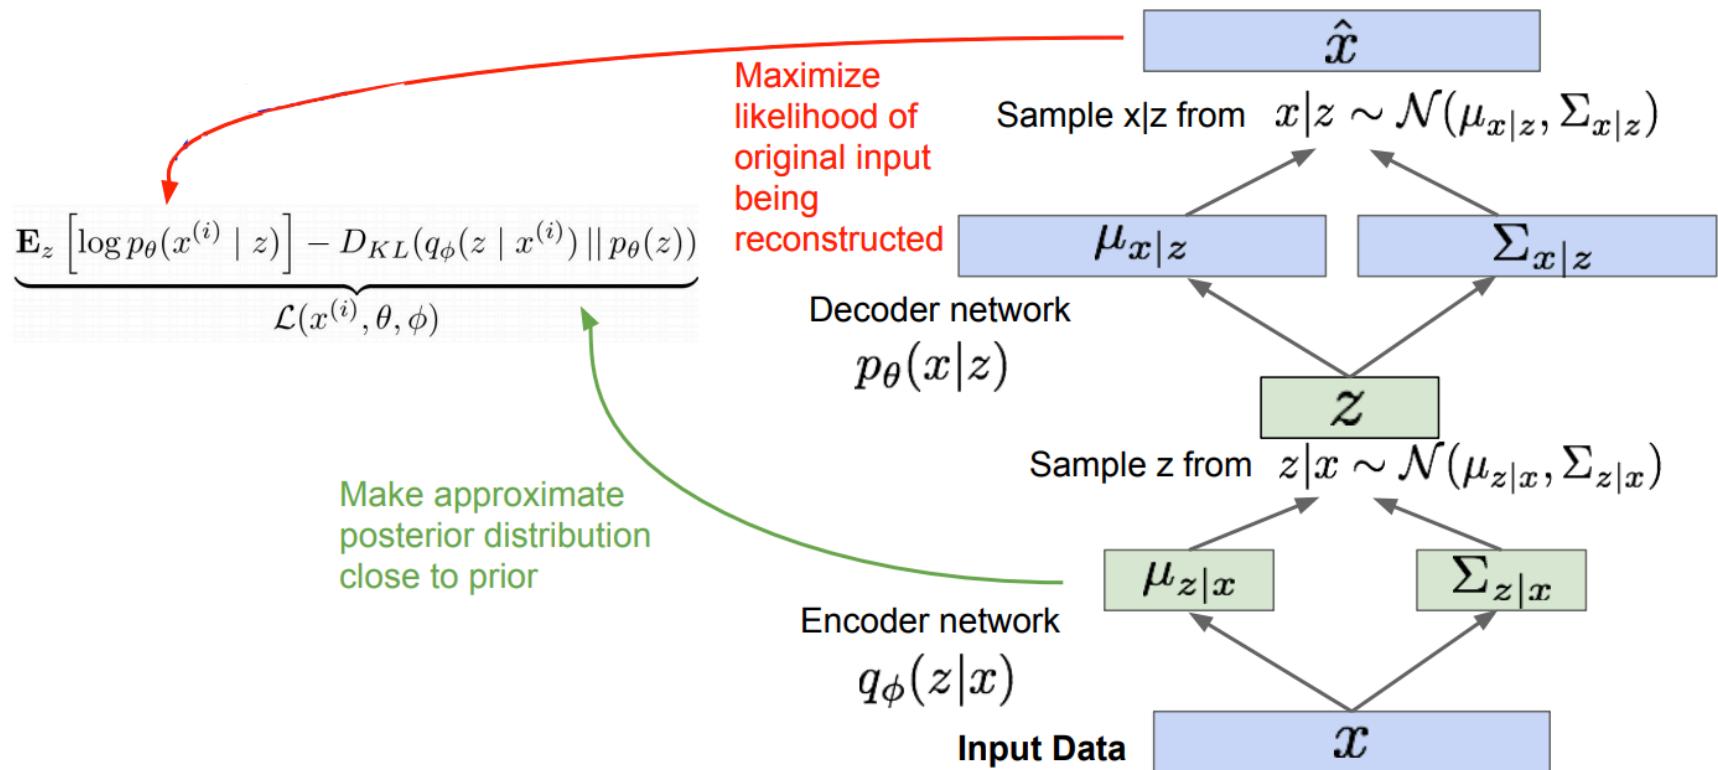
$$\begin{aligned}\log p_\theta(x^{(i)}) &= \mathbf{E}_{z \sim q_\phi(z|x^{(i)})} [\log p_\theta(x^{(i)})] \quad (p_\theta(x^{(i)}) \text{ Does not depend on } z) \\ &= \mathbf{E}_z \left[ \log \frac{p_\theta(x^{(i)} | z)p_\theta(z)}{p_\theta(z | x^{(i)})} \right] \quad (\text{Bayes' Rule}) \\ &= \mathbf{E}_z \left[ \log \frac{p_\theta(x^{(i)} | z)p_\theta(z)}{p_\theta(z | x^{(i)})} \frac{q_\phi(z | x^{(i)})}{q_\phi(z | x^{(i)})} \right] \quad (\text{Multiply by constant}) \\ &= \mathbf{E}_z [\log p_\theta(x^{(i)} | z)] - \mathbf{E}_z \left[ \log \frac{q_\phi(z | x^{(i)})}{p_\theta(z)} \right] + \mathbf{E}_z \left[ \log \frac{q_\phi(z | x^{(i)})}{p_\theta(z | x^{(i)})} \right] \quad (\text{Logarithms}) \\ &= \underbrace{\mathbf{E}_z [\log p_\theta(x^{(i)} | z)] - D_{KL}(q_\phi(z | x^{(i)}) || p_\theta(z))}_{\mathcal{L}(x^{(i)}, \theta, \phi)} + \underbrace{D_{KL}(q_\phi(z | x^{(i)}) || p_\theta(z | x^{(i)}))}_{\geq 0}\end{aligned}$$

Tractable lower bound which we can take gradient of and optimize! ( $p_\theta(x|z)$  differentiable, KL term differentiable)

intractable

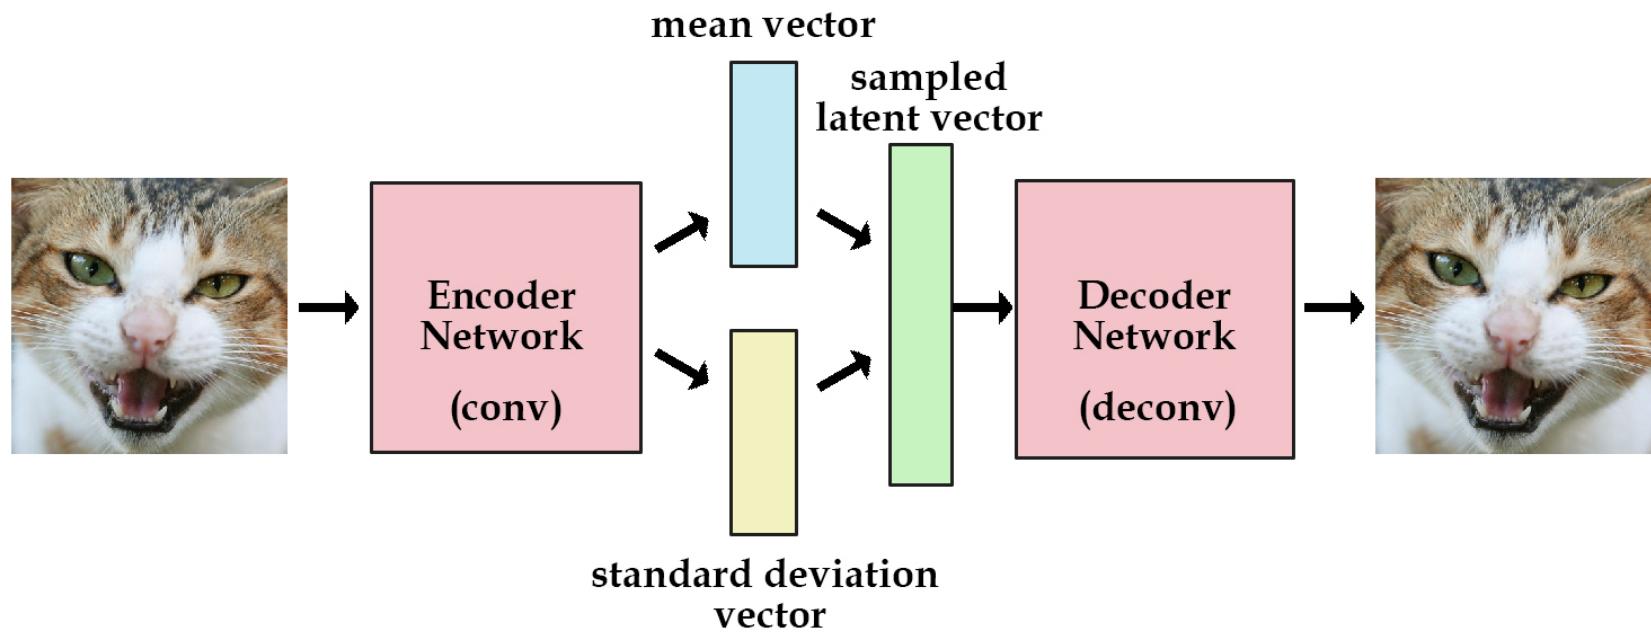


# Recap: Variational Autoencoders





# Recap: Variational Autoencoders





# Recap: Variational Autoencoders

Defines an intractable density

=> derive and optimize a (vairational) lower bound

Probabilistic spin to traditional autoencoders

=> allows generating data



# Recap: Variational Autoencoders

## Pros:

- Principled approach to generate models
- Allows inference of  $q(z|x)$ , can be useful feature representation for other tasks

## Cons:

- Maximizes lower bound of likelihood: okay, but not as good evaluation as tractable density
- Samples blurrier and lower quality compared to state-of-the-art (GANs)

## Active areas of research:

- More flexible approximations, e.g. richer approximate posterior instead of diagonal Gaussian
- Incorporating structure in latent variables

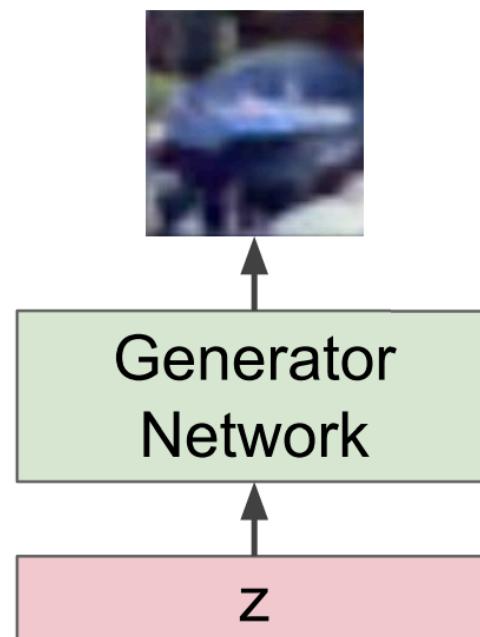


# Recap: Generative Adversarial Networks

Try a new way

Output: Sample from  
training distribution

Input: Random noise



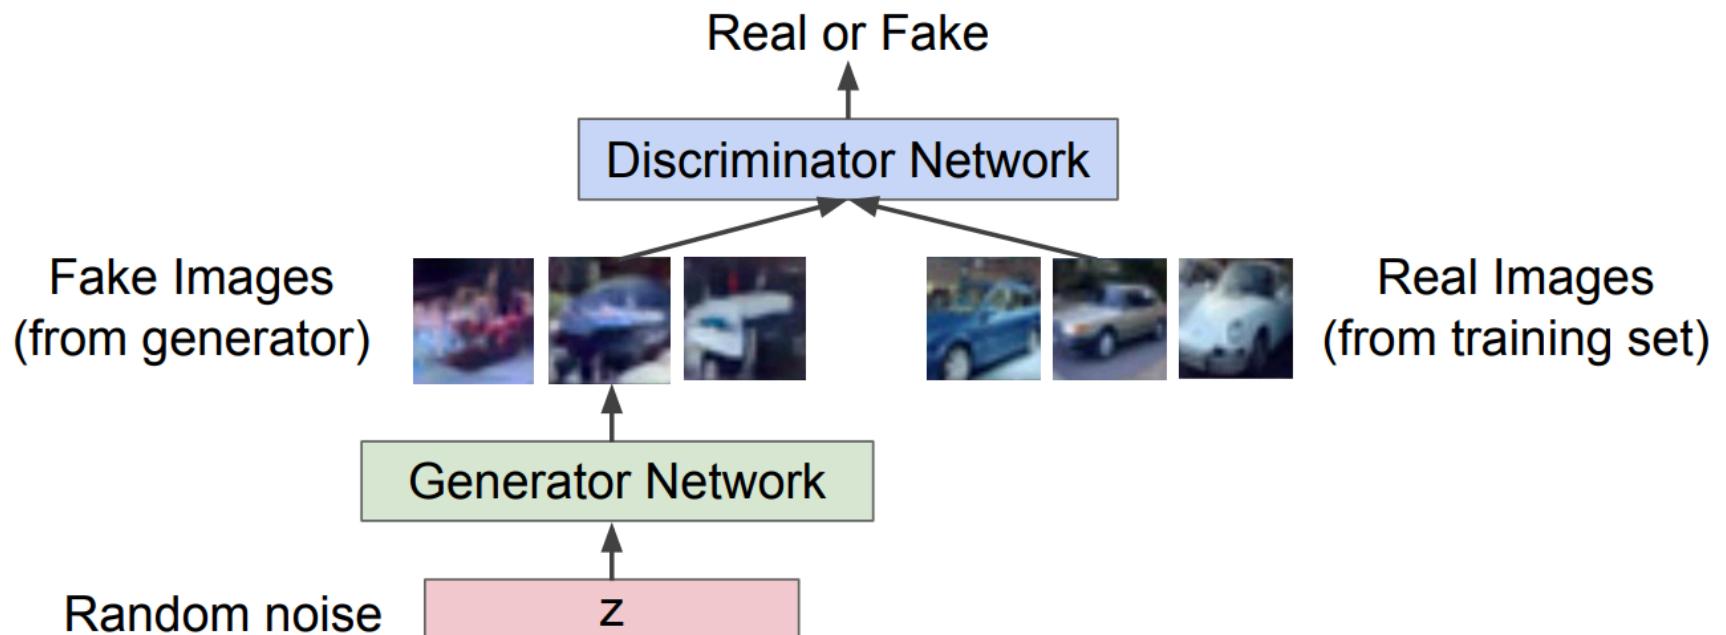


# Recap: Generative Adversarial Networks

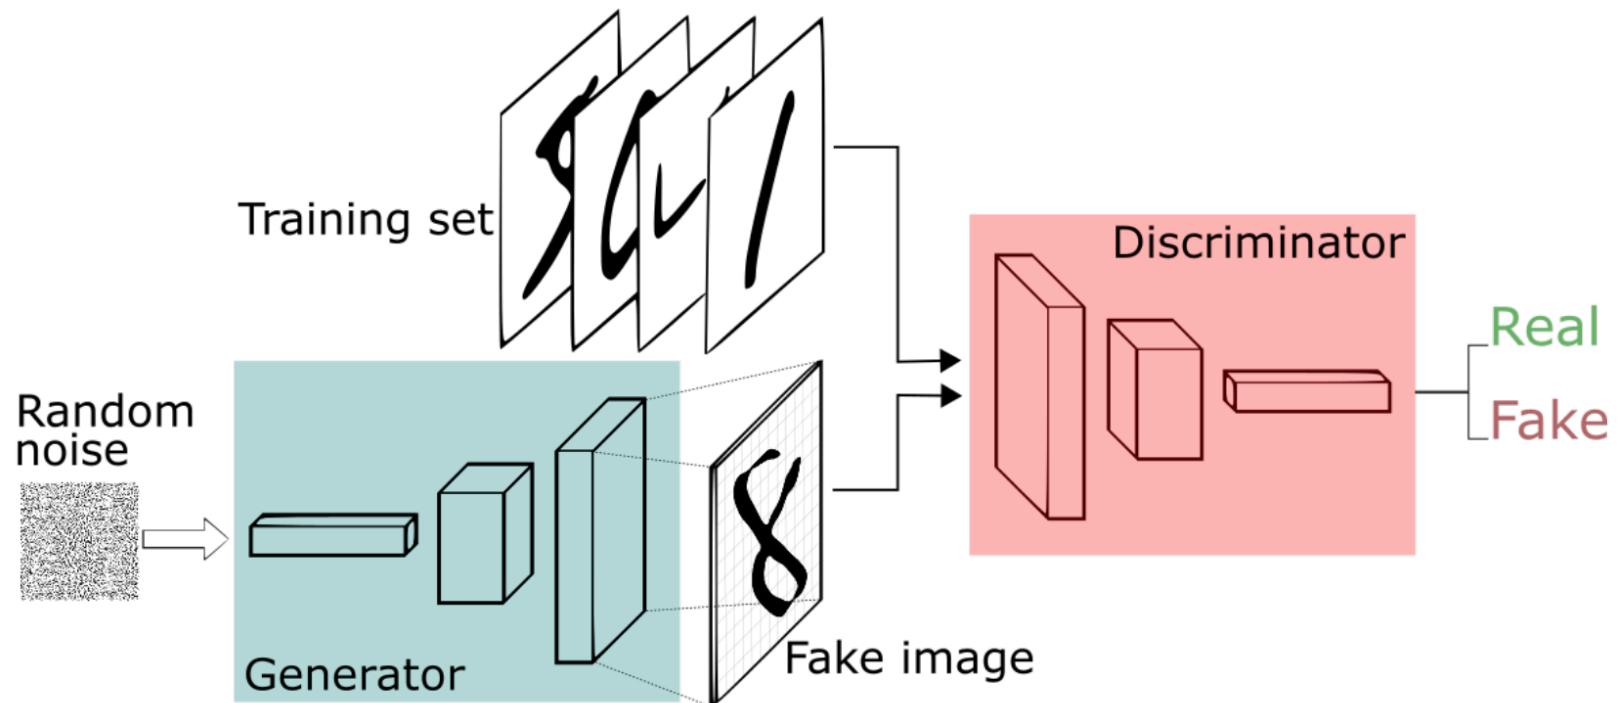
## Architecture

**Generator network:** try to fool the discriminator by generating real-looking images

**Discriminator network:** try to distinguish between real and fake images



# Recap: Generative Adversarial Networks





# Recap: Generative Adversarial Networks

Training process: two-player game

Train jointly in **minimax game**

Minimax objective function:

$$\min_{\theta_g} \max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log \underbrace{D_{\theta_d}(x)}_{\text{Discriminator output for real data } x} + \mathbb{E}_{z \sim p(z)} \log(1 - \underbrace{D_{\theta_d}(G_{\theta_g}(z))}_{\text{Discriminator output for generated fake data } G(z)}) \right]$$

Discriminator outputs likelihood in (0,1) of real image

Discriminator ( $\theta_d$ ) wants to maximize objective such that  $D(x)$  is close to 1 (real) and  $D(G(z))$  is close to 0 (fake)

Generator ( $\theta_g$ ) wants to minimize objective such that  $D(G(z))$  is close to 1 (discriminator is fooled into thinking generated  $G(z)$  is real)



# Recap: Generative Adversarial Networks

## Training procedure

Use SGD-like algorithm of choice (Adam) on two mini-batches simultaneously:

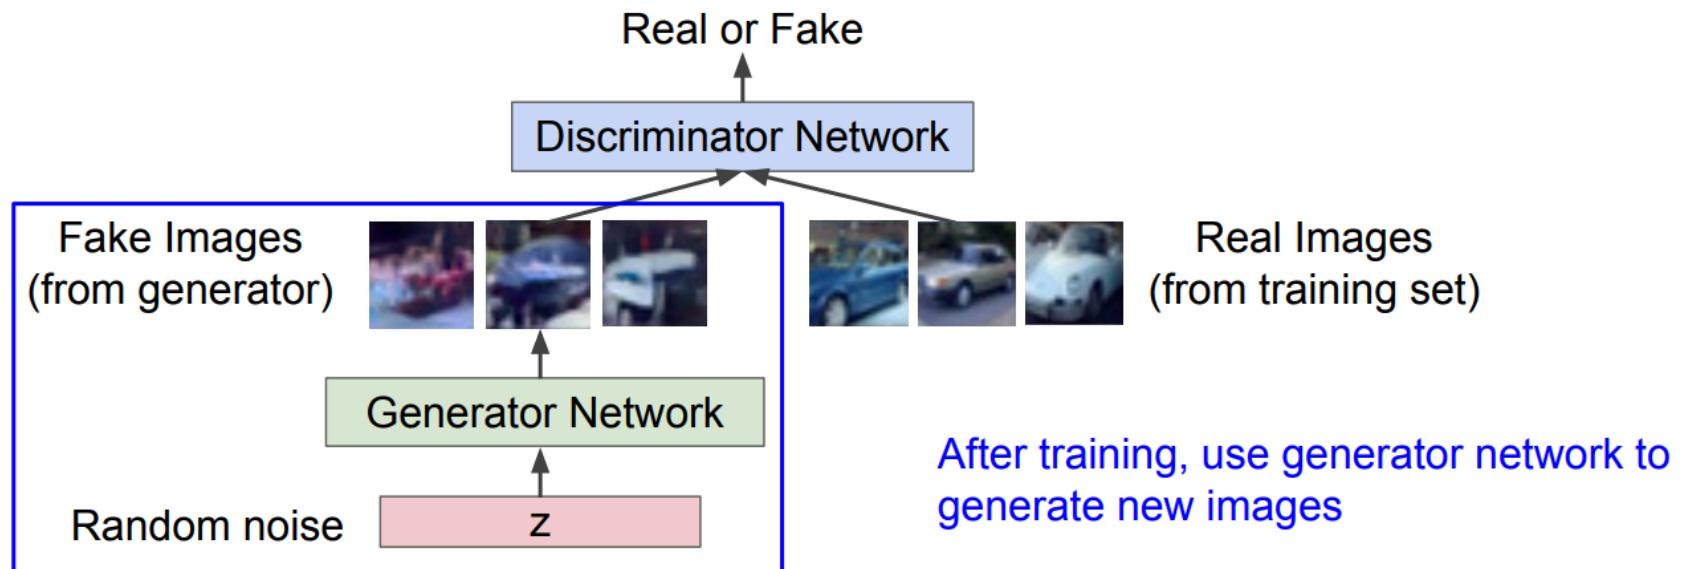
- A mini-batch of training examples
- A mini-batch of generated examples

Optional: run  $k$  steps of one player for every step of the other player



# Recap: Generative Adversarial Networks

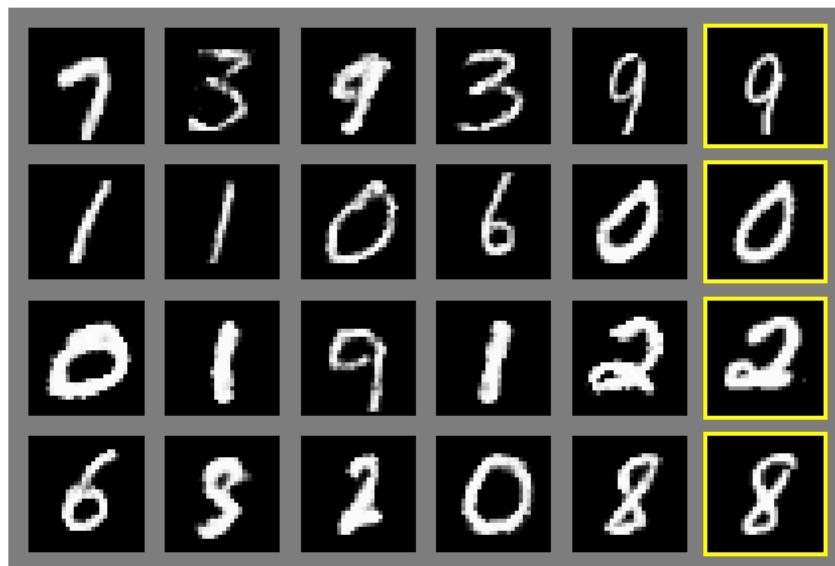
Generate new images





# Recap: Generative Adversarial Networks

Generated samples



MNIST



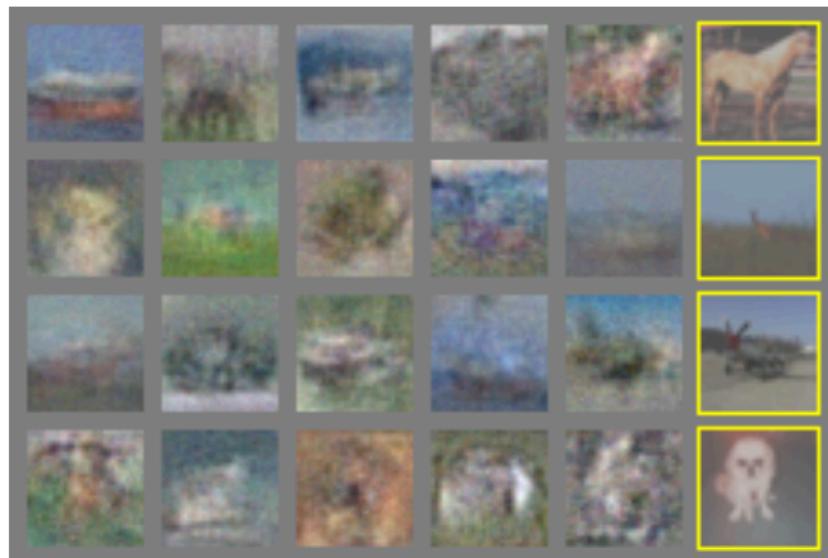
TFD

Nearest neighbor from training set

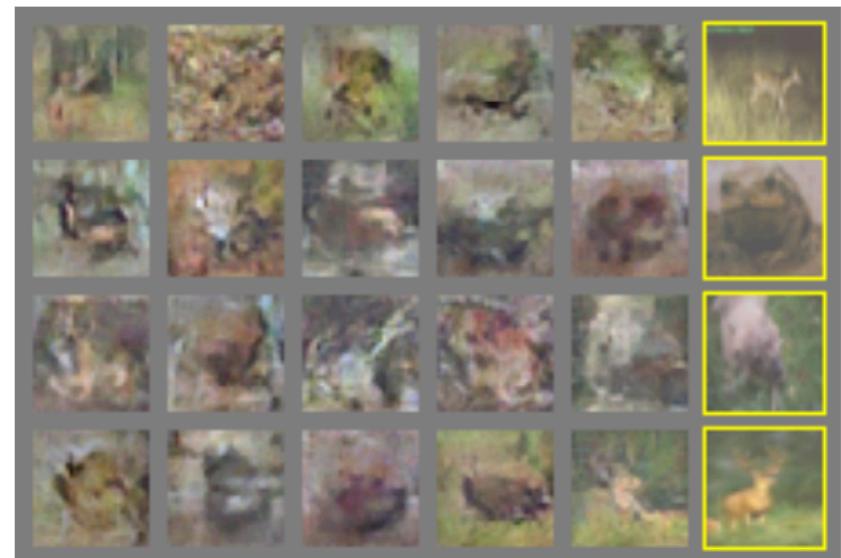


# Recap: Generative Adversarial Networks

Generated samples



CIFAR-10 (FC)



CIFAR-10 (CNN)

Nearest neighbor from training set



# Recap: Generative Adversarial Networks

What is the latent space like?



You can interpolate along the hidden space to produce smooth transitions of images.



# Recap: Key differences between VAEs and GANs

- VAEs are more theoretically grounded than GANs. GANs are more based on what works.
- GANs traditionally only learn the decoder (generator) but there are variations that learn an encoder as well; there are some problems where you want both and some problems where just the decoder will suffice. VAEs learn an encoder/decoder pair.
- GAN decoder sees samples from prior  $q(z)$ , VAE decoder sees samples from model  $p(z | x)$ .



# Recap: Most important difference between VAEs and GANs

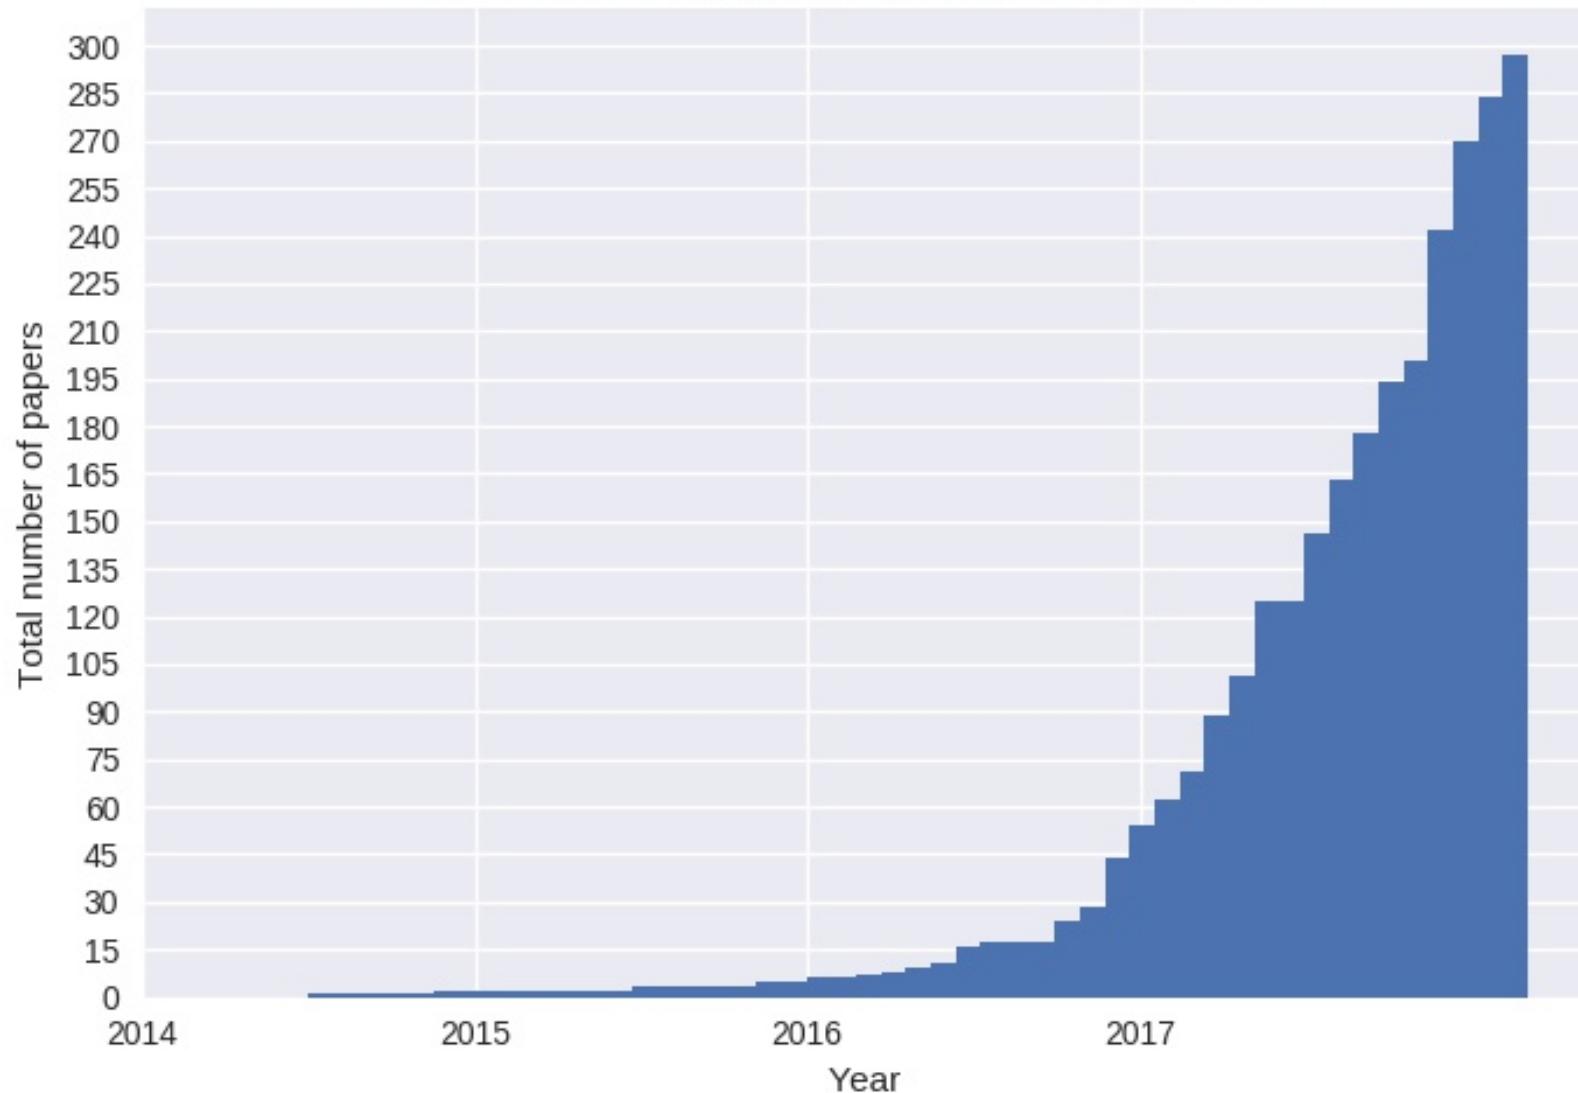
“Distance”:

- VAE objective for the decoder is some **man-made** objective function, like L2 distance between images
- GAN objective for the generator is some complicated objective function **defined by a neural network**



# Generative Adversarial Networks

Cumulative number of named GAN papers by month



2017: Year of GAN



# Overview

- Some GANs
- Optimization issues of GANs
- Optimization techniques of GANs
- Applications of GANs
- TensorFlow tutorial

**Let's see some GANs!**



# Some Generative Adversarial Networks

## Conditional GAN (CGAN)

CGAN includes a label and learn  $p(x|y)$

Generator learns  $p(x|z, y)$

Discriminator learns  $p(\text{label}|x, y)$

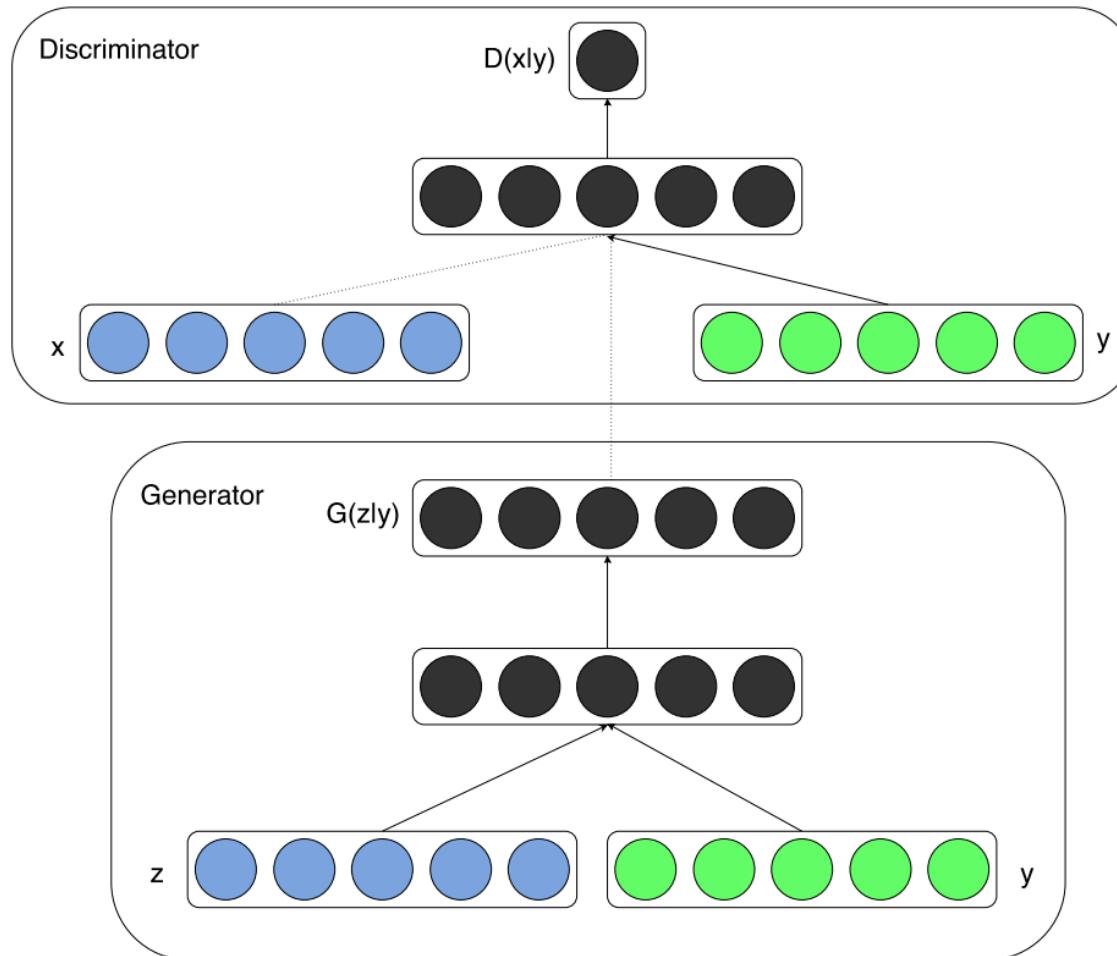
GAN loss:  $\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})}[\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})}[\log(1 - D(G(\mathbf{z})))].$

CGAN loss:  $\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})}[\log D(\mathbf{x}|\mathbf{y})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})}[\log(1 - D(G(\mathbf{z}|\mathbf{y})))].$



# Some Generative Adversarial Networks

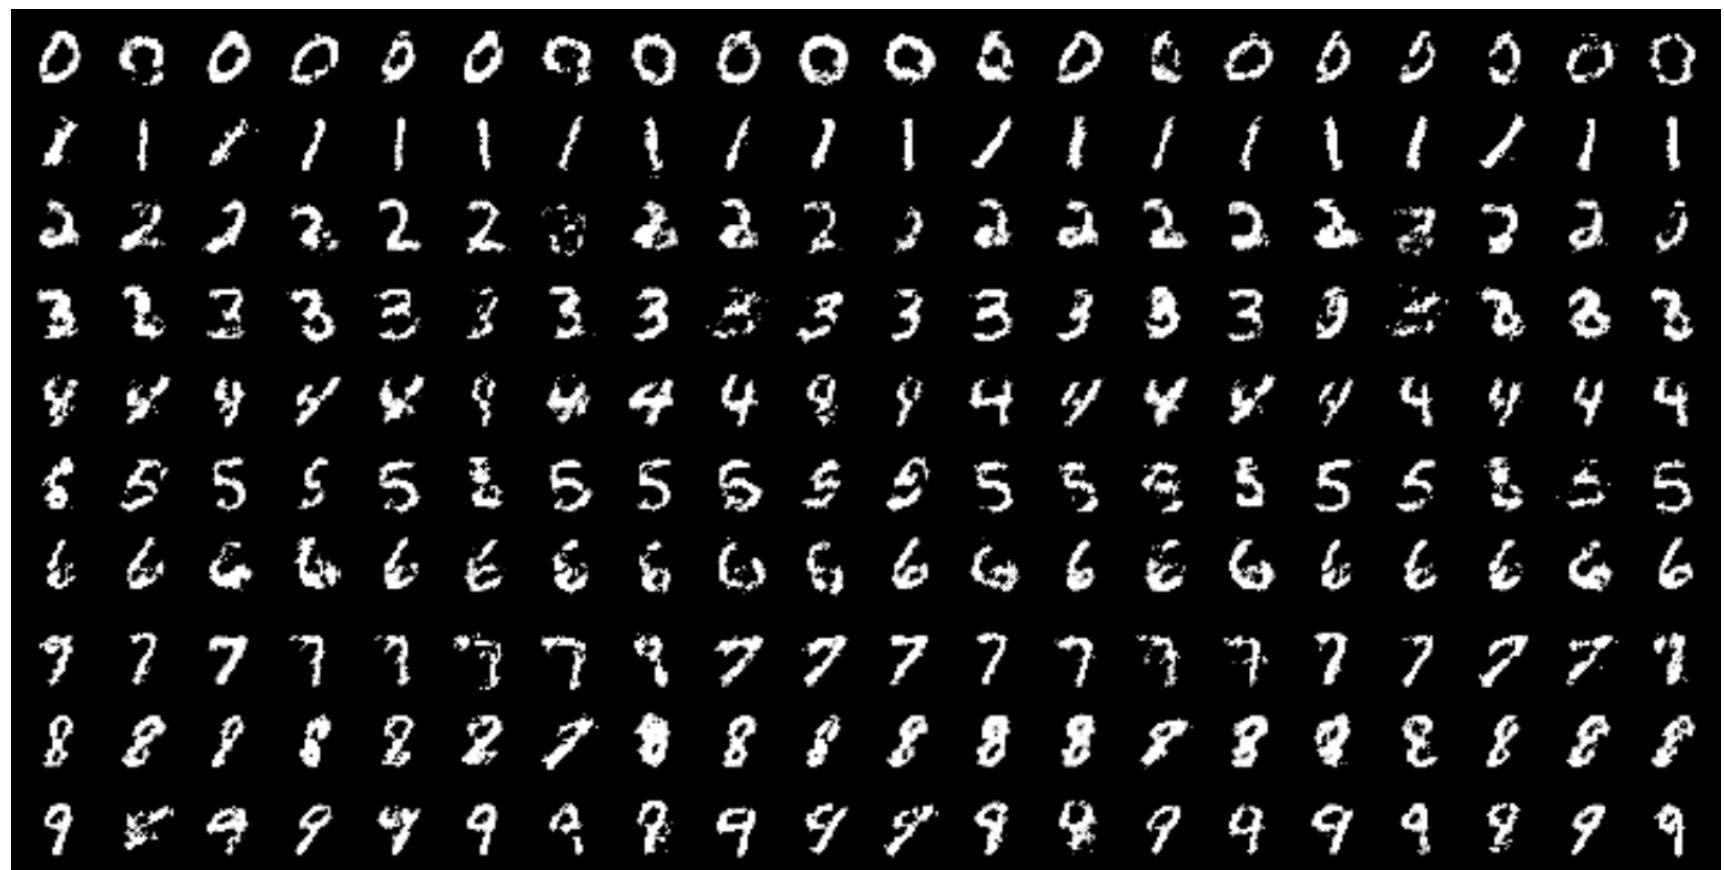
## Conditional GAN (CGAN)





# Some Generative Adversarial Networks

## Conditional GAN (CGAN)





# Some Generative Adversarial Networks

## DCGAN

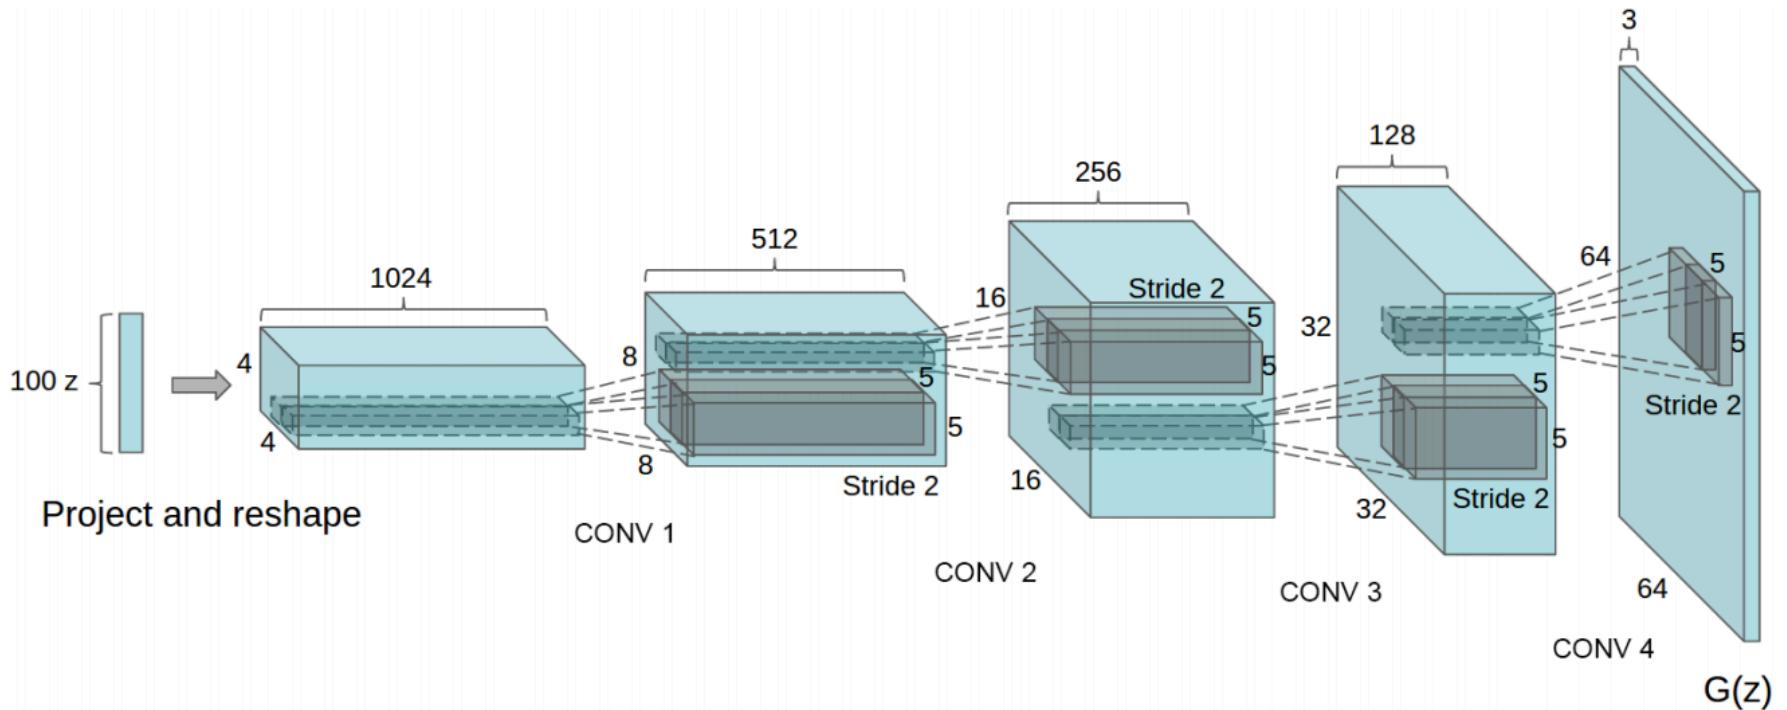
DCGAN made several improvements allowing GANs to be trained on larger/deeper CNNs

- Using **Leaky ReLUs** in discriminator for all layers
- Use **ReLU** in generator for all layers except for the output, which uses **Tanh**
- Using **BatchNorm** in both the generator and the discriminator
- Replace any pooling layers with **strided convolutions**
- Remove fully connected hidden layers for deeper architectures



# Some Generative Adversarial Networks

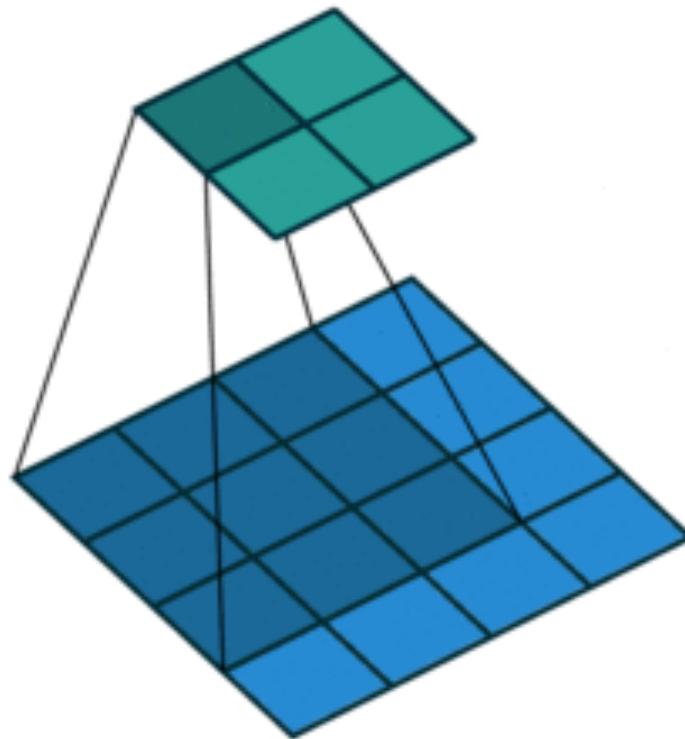
## DCGAN generator





# Some Generative Adversarial Networks

## Convolution



(No padding, unit strides) Convolving a  $3 \times 3$  kernel over a  $4 \times 4$  input using unit strides (i.e.,  $i = 4$ ,  $k = 3$ ,  $s = 1$  and  $p = 0$ ).



# Some Generative Adversarial Networks

## Convolution

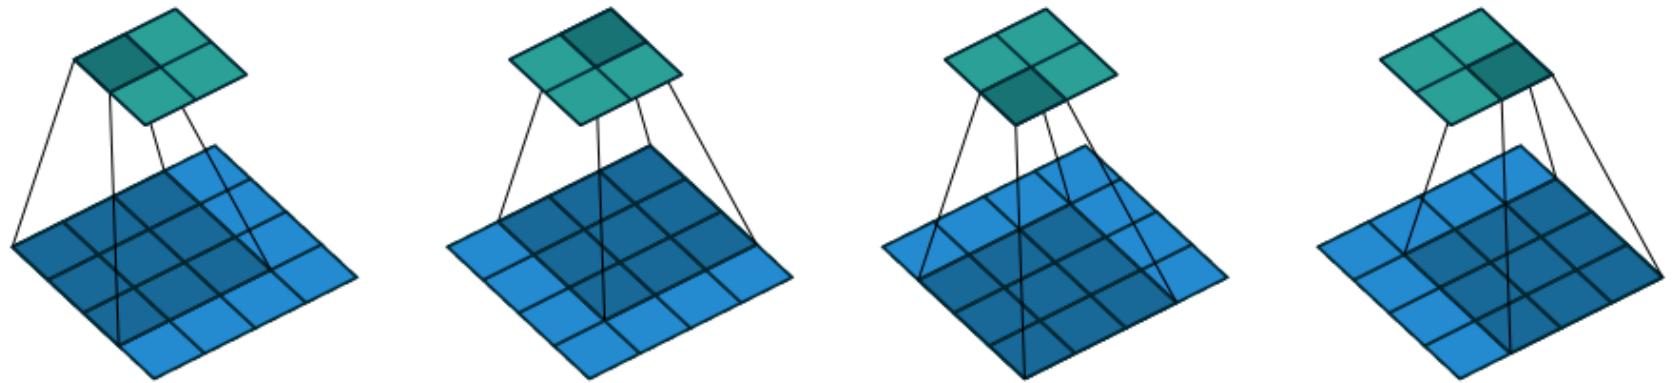


Figure 2.1: (No padding, unit strides) Convolving a  $3 \times 3$  kernel over a  $4 \times 4$  input using unit strides (i.e.,  $i = 4$ ,  $k = 3$ ,  $s = 1$  and  $p = 0$ ).



# Some Generative Adversarial Networks

## Convolution as a matrix operation

$$\begin{pmatrix} w_{0,0} & w_{0,1} & w_{0,2} & 0 & w_{1,0} & w_{1,1} & w_{1,2} & 0 & w_{2,0} & w_{2,1} & w_{2,2} & 0 & 0 & 0 & 0 & 0 \\ 0 & w_{0,0} & w_{0,1} & w_{0,2} & 0 & w_{1,0} & w_{1,1} & w_{1,2} & 0 & w_{2,0} & w_{2,1} & w_{2,2} & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & w_{0,0} & w_{0,1} & w_{0,2} & 0 & w_{1,0} & w_{1,1} & w_{1,2} & 0 & w_{2,0} & w_{2,1} & w_{2,2} & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & w_{0,0} & w_{0,1} & w_{0,2} & 0 & w_{1,0} & w_{1,1} & w_{1,2} & 0 & w_{2,0} & w_{2,1} & w_{2,2} & 0 \end{pmatrix}$$

×

unrolled input vector (16×1)

||

unrolled output vector (4×1)



# Some Generative Adversarial Networks

## Convolution as a matrix operation

$$\begin{pmatrix} w_{0,0} & w_{0,1} & w_{0,2} & 0 & w_{1,0} & w_{1,1} & w_{1,2} & 0 & w_{2,0} & w_{2,1} & w_{2,2} & 0 & 0 & 0 & 0 & 0 \\ 0 & w_{0,0} & w_{0,1} & w_{0,2} & 0 & w_{1,0} & w_{1,1} & w_{1,2} & 0 & w_{2,0} & w_{2,1} & w_{2,2} & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & w_{0,0} & w_{0,1} & w_{0,2} & 0 & w_{1,0} & w_{1,1} & w_{1,2} & 0 & w_{2,0} & w_{2,1} & w_{2,2} & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & w_{0,0} & w_{0,1} & w_{0,2} & 0 & w_{1,0} & w_{1,1} & w_{1,2} & 0 & w_{2,0} & w_{2,1} & w_{2,2} & 0 \end{pmatrix}$$

$C$

$\times$

$V_{high}$  unrolled input vector ( $16 \times 1$ )

$\parallel$

$V_{low}$  unrolled output vector ( $4 \times 1$ )



# Some Generative Adversarial Networks

Convolution as a matrix operation

$$C \times V_{high} = V_{low}$$



# Some Generative Adversarial Networks

Convolution as a matrix operation

$$C \times V_{high} = V_{low}$$

$$V_{low} \xrightarrow{?} V_{high}$$



# Some Generative Adversarial Networks

Convolution vs Transposed convolution

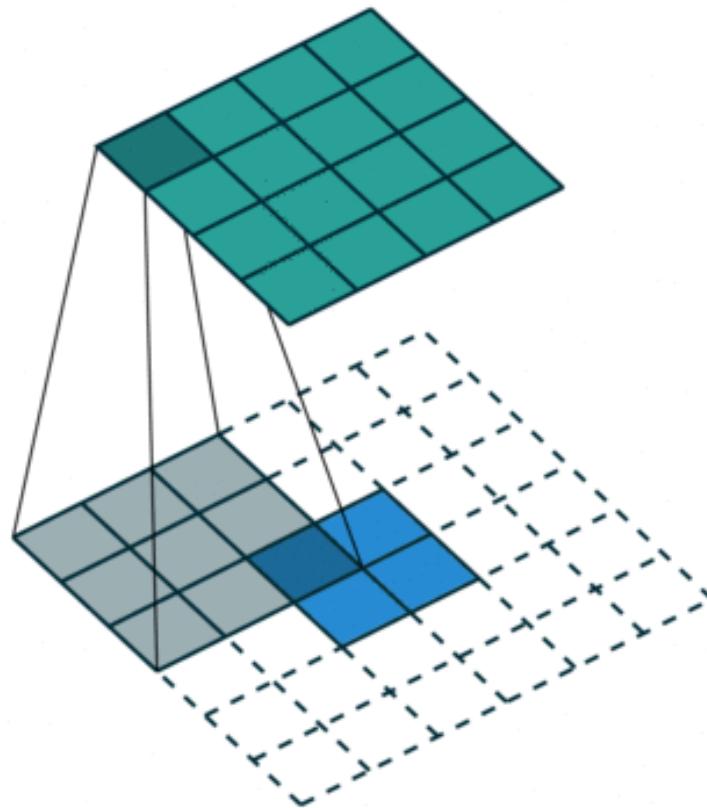
$$C \times V_{high} = V_{low}$$

$$V_{high} = C^T \times V_{low}$$



# Some Generative Adversarial Networks

## Transposed convolution





# Some Generative Adversarial Networks

## Transposed convolution

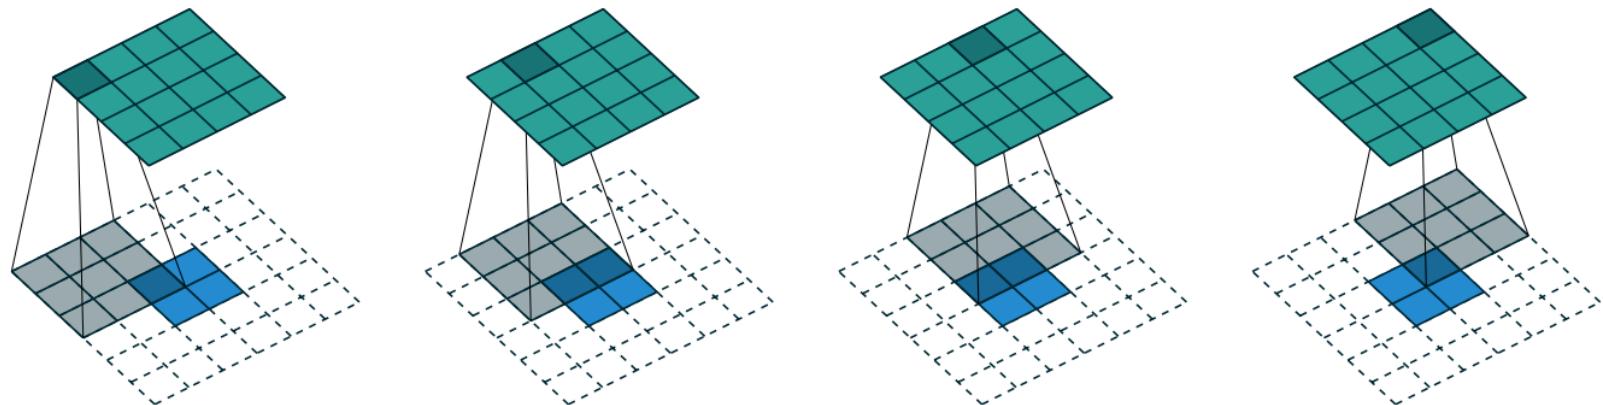
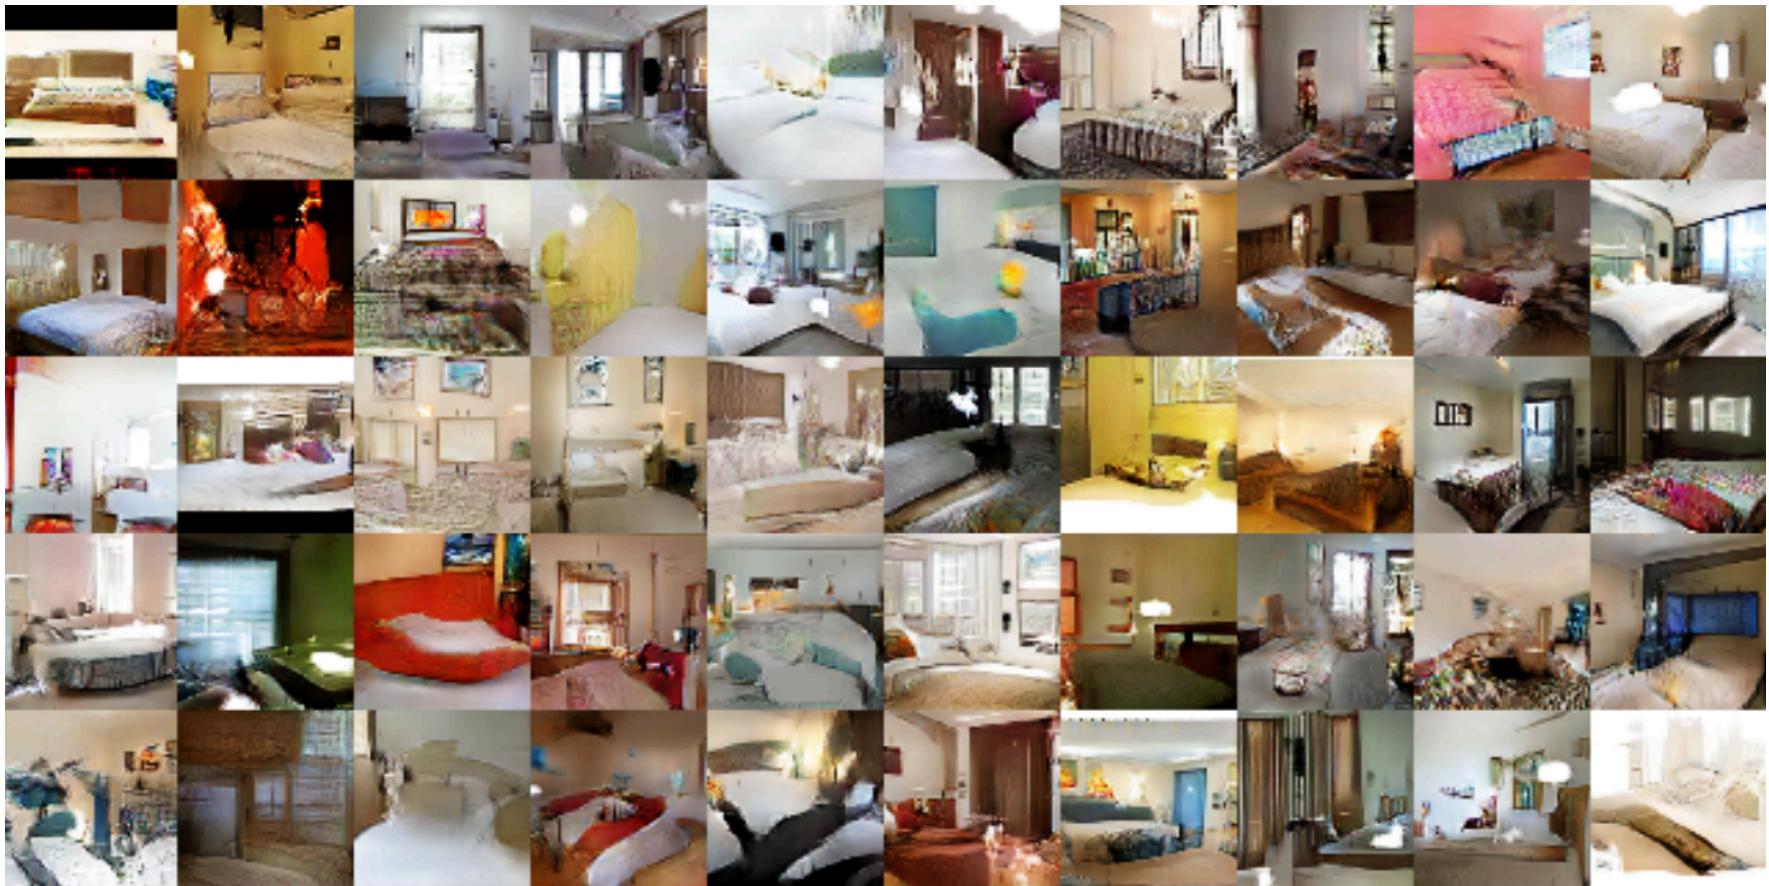


Figure 4.1: The transpose of convolving a  $3 \times 3$  kernel over a  $4 \times 4$  input using unit strides (i.e.,  $i = 4$ ,  $k = 3$ ,  $s = 1$  and  $p = 0$ ). It is equivalent to convolving a  $3 \times 3$  kernel over a  $2 \times 2$  input padded with a  $2 \times 2$  border of zeros using unit strides (i.e.,  $i' = 2$ ,  $k' = k$ ,  $s' = 1$  and  $p' = 2$ ).



# Some Generative Adversarial Networks

## DCGAN results





# Some Generative Adversarial Networks

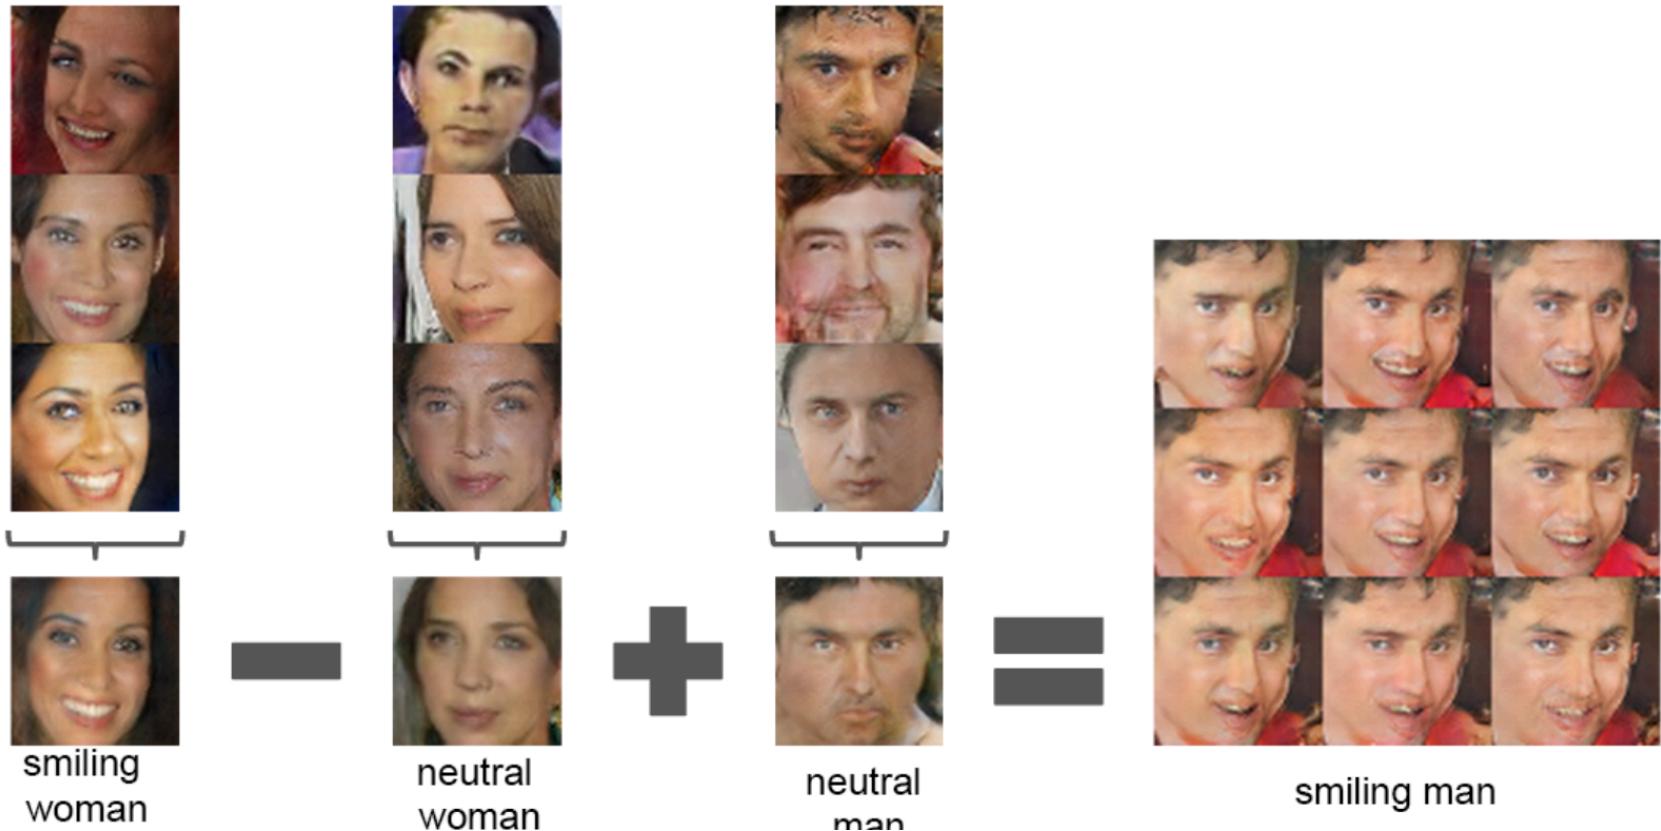
## DCGAN results





# Some Generative Adversarial Networks

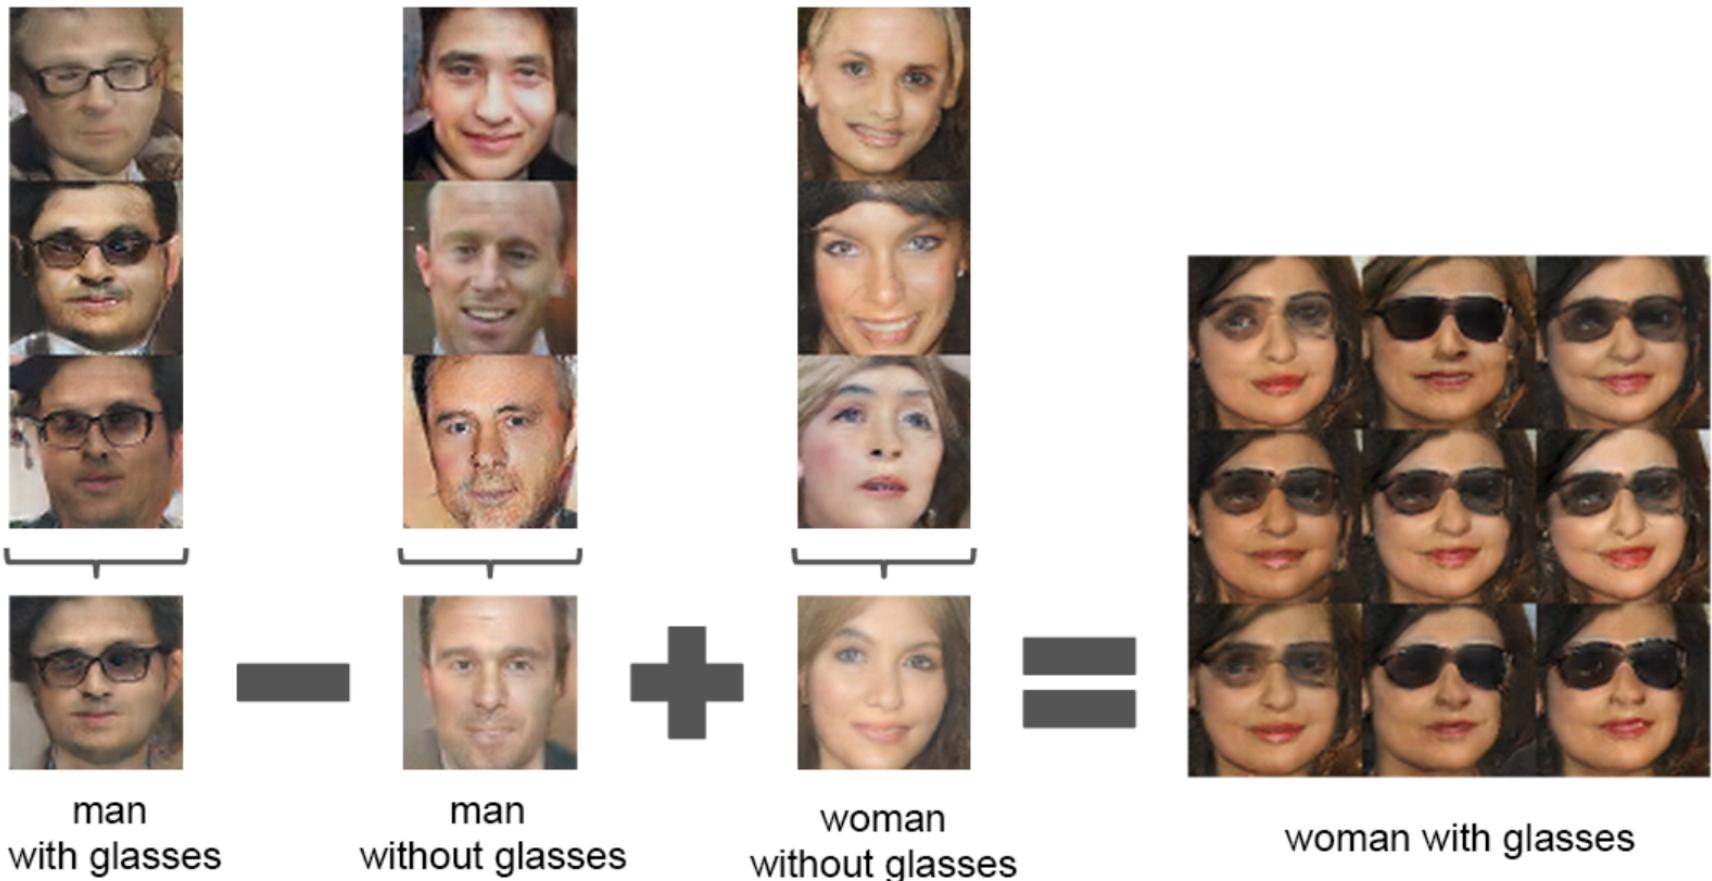
## DCGAN vector math





# Some Generative Adversarial Networks

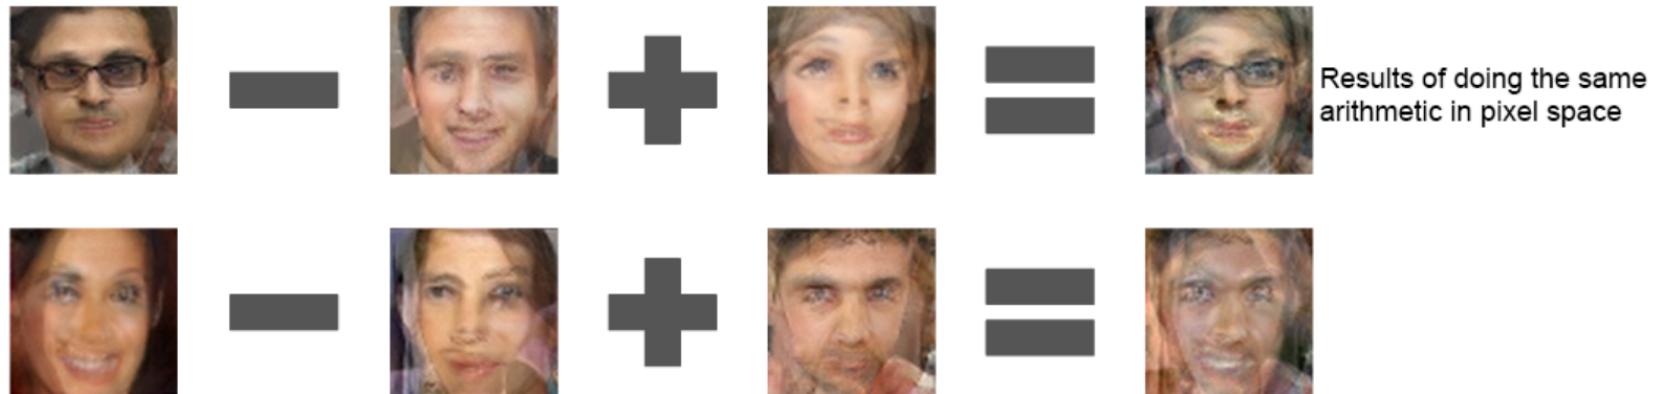
## DCGAN





# Some Generative Adversarial Networks

## DCGAN vector math





# Some Generative Adversarial Networks

DCGAN vector manipulation





# Some Generative Adversarial Networks

DCGAN online demos

Neural Face: <https://carpedm20.github.io/faces/>

DCGAN face generator: <http://mattyai.github.io/chainer-DCGAN/>

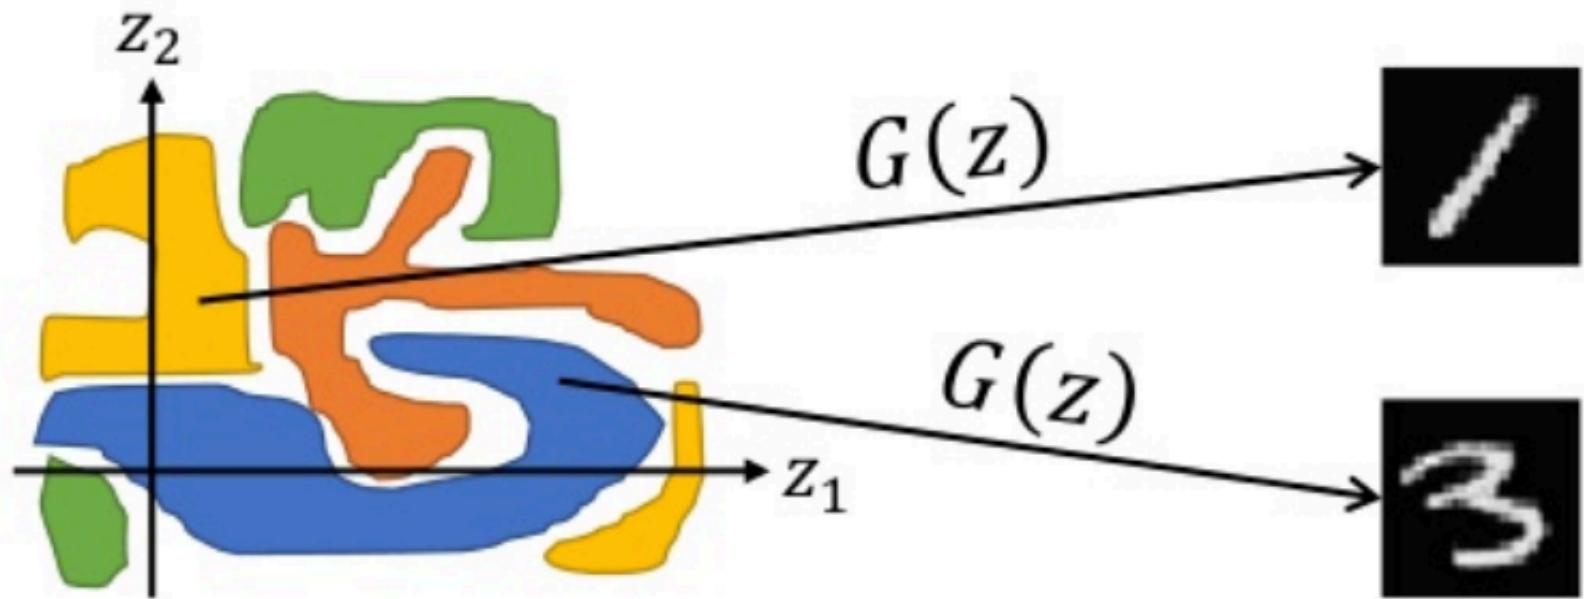
MakeGirlsMoe: <https://make.girls.moe/#/>



# Some Generative Adversarial Networks

## InfoGAN

Highly entangled noise

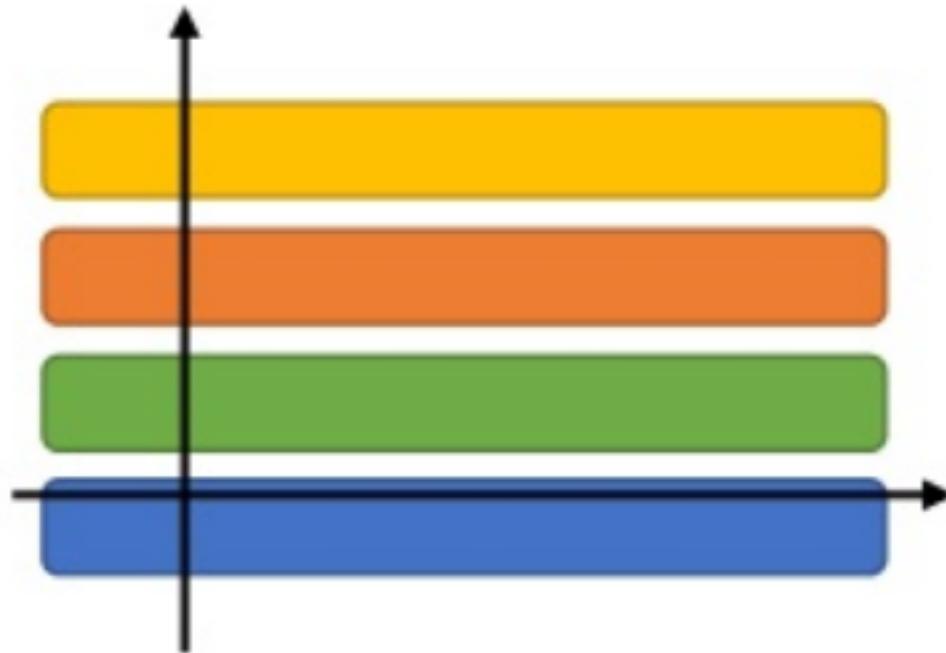




# Some Generative Adversarial Networks

## InfoGAN

Disentangled noise





# Some Generative Adversarial Networks

InfoGAN

Information theory

Mutual information:

$$I(X; Y) = H(X) - H(X|Y) = H(Y) - H(Y|X)$$

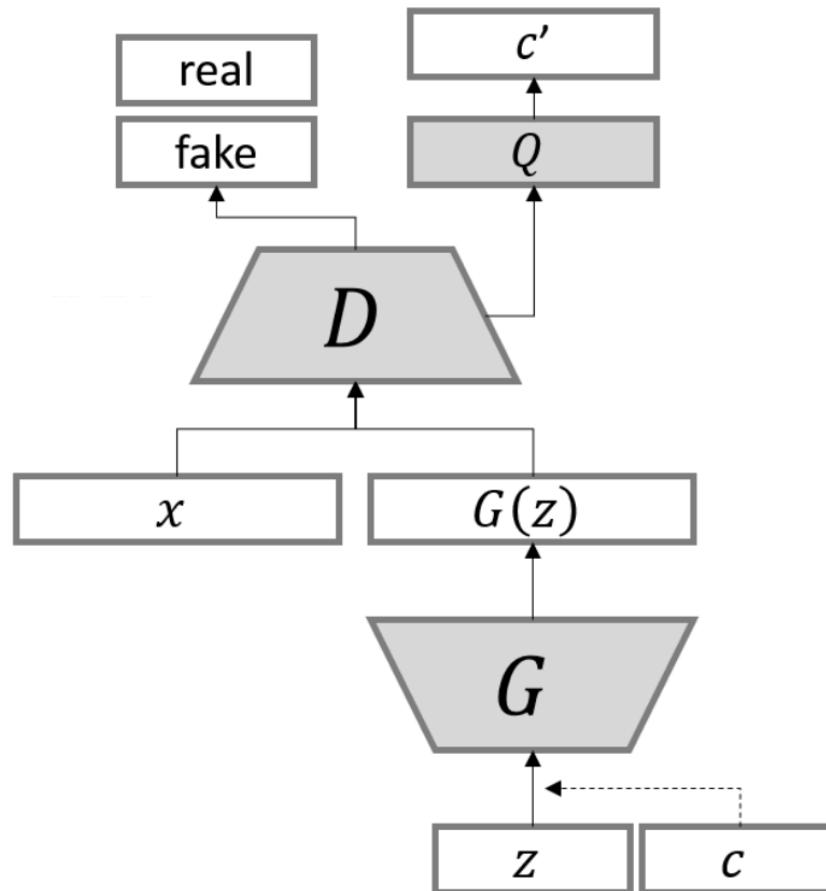
Objective function:

$$\min_G \max_D V_I(D, G) = V(D, G) - \lambda I(c; G(z, c))$$



# Some Generative Adversarial Networks

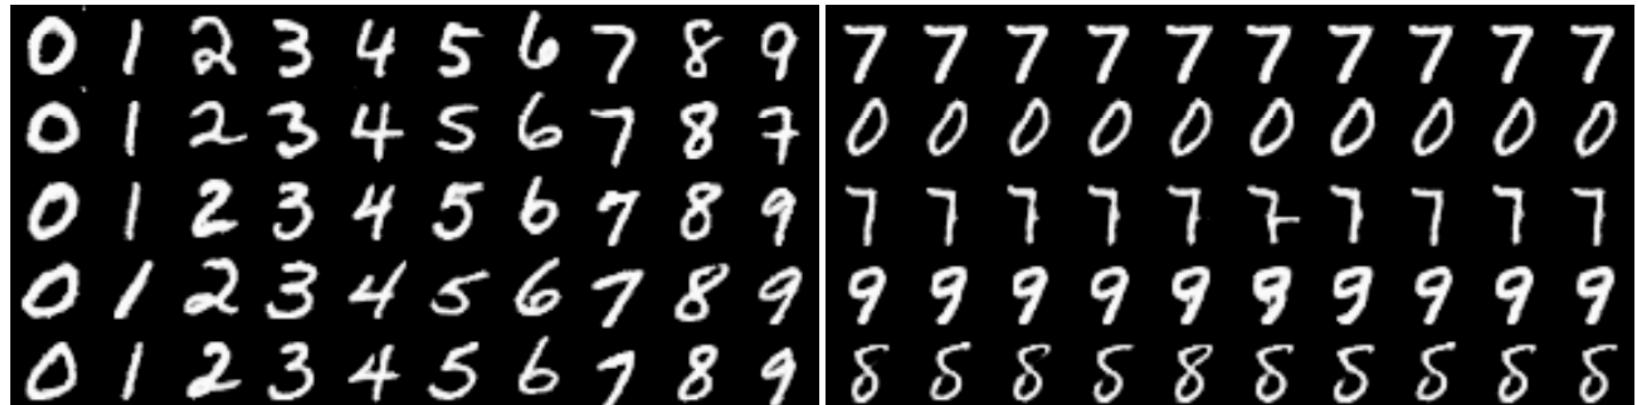
## InfoGAN architecture





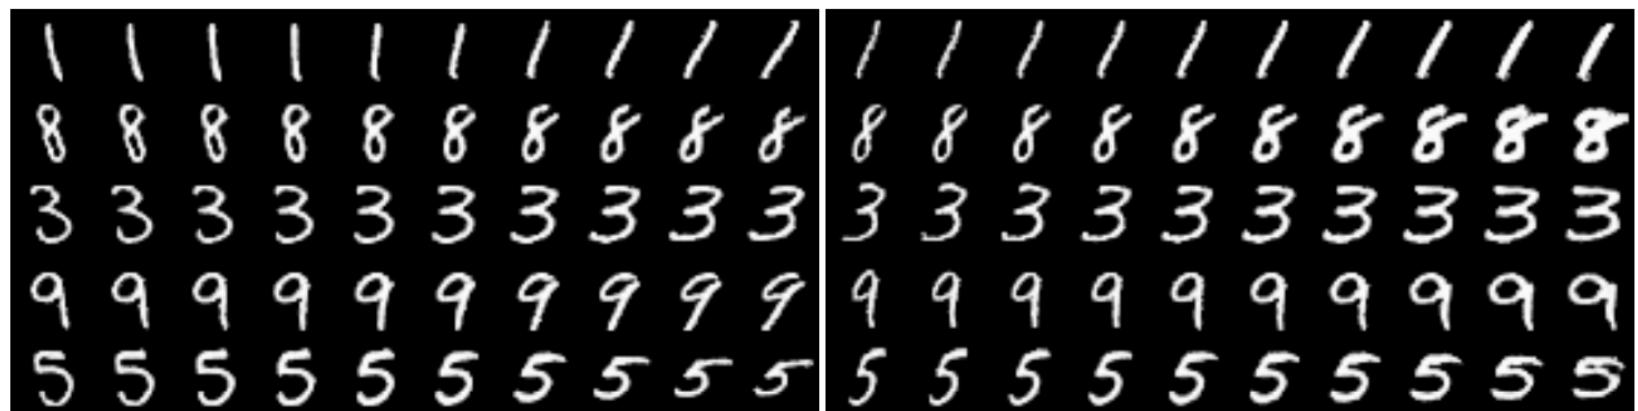
# Some Generative Adversarial Networks

## InfoGAN results



(a) Varying  $c_1$  on InfoGAN (Digit type)

(b) Varying  $c_1$  on regular GAN (No clear meaning)



(c) Varying  $c_2$  from  $-2$  to  $2$  on InfoGAN (Rotation)

(d) Varying  $c_3$  from  $-2$  to  $2$  on InfoGAN (Width)



# Some Generative Adversarial Networks

## InfoGAN results



(a) Azimuth (pose)

(b) Elevation



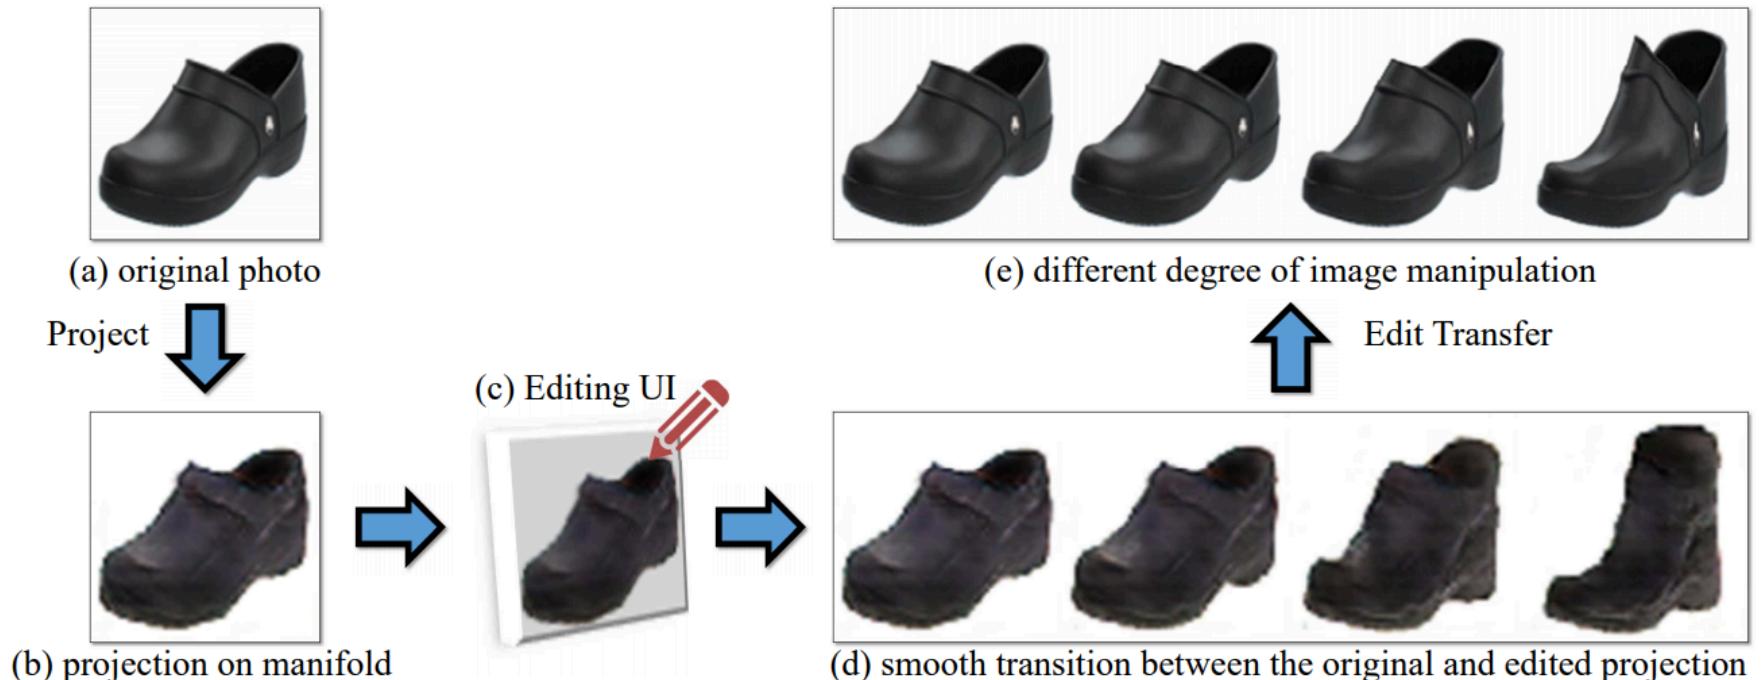
(c) Lighting

(d) Wide or Narrow



# Some Generative Adversarial Networks

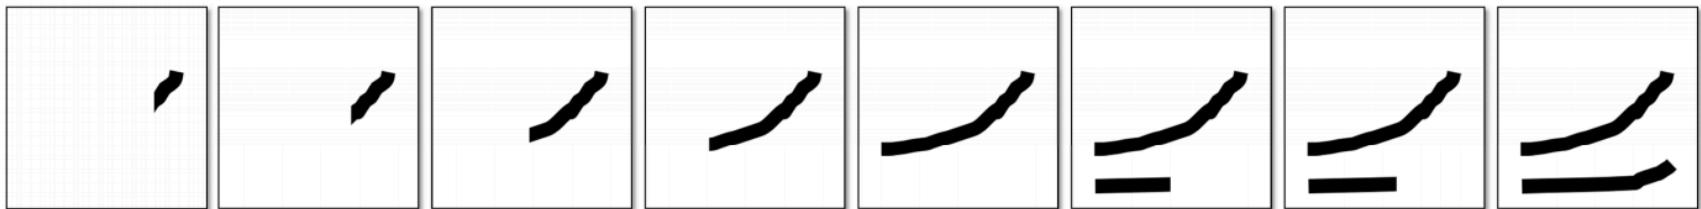
iGAN





# Some Generative Adversarial Networks

iGAN



(a) User constraints  $v_g$  at different update steps



$G(z_0)$

(b) Updated images according to user edits

$G(z_1)$



(c) Linear interpolation between  $G(z_0)$  and  $G(z_1)$

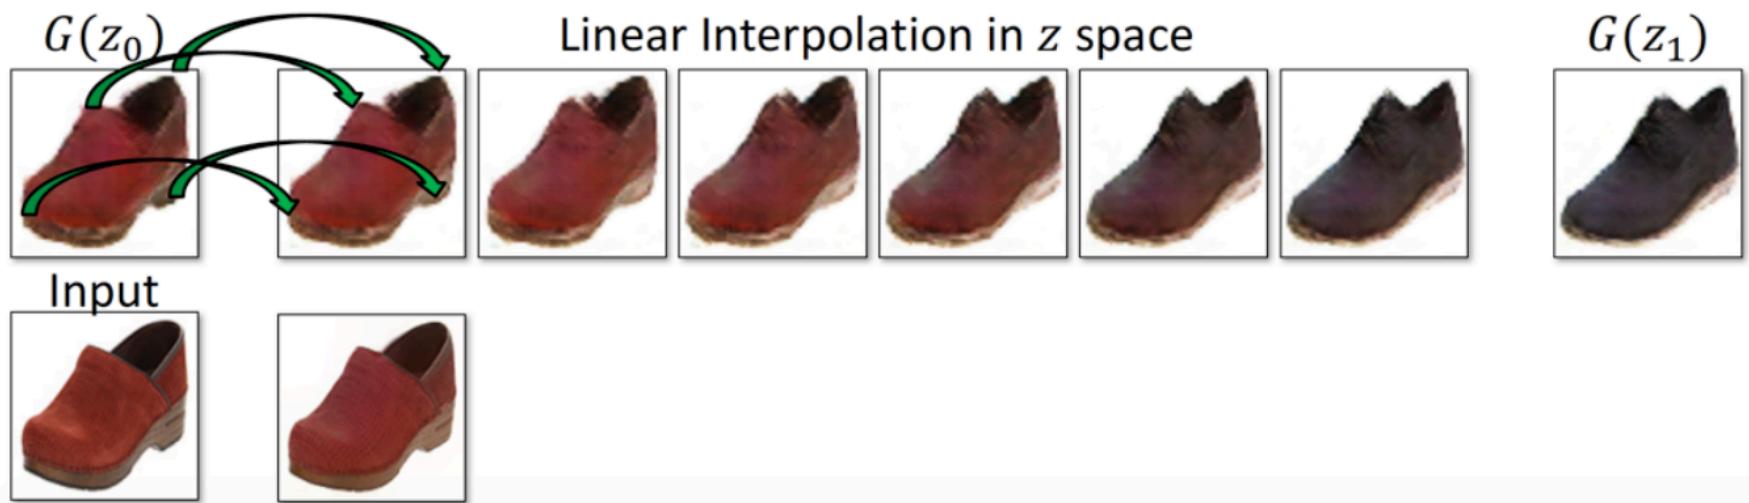


# Some Generative Adversarial Networks

## iGAN

**Motion ( $u, v$ ) + Color ( $A_{3 \times 4}$ ):** estimate per-pixel geometric and color variation

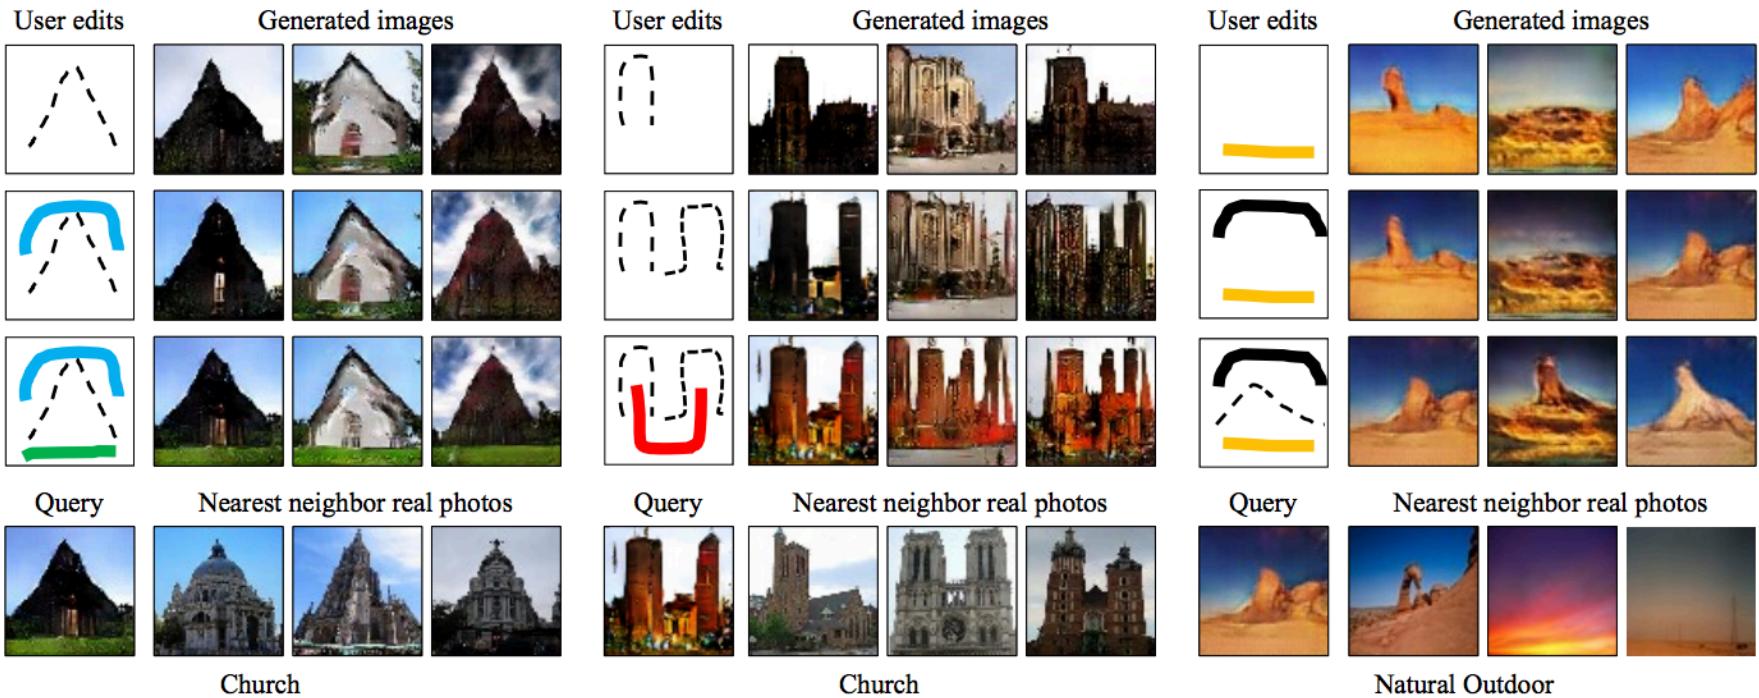
$$\iint \underbrace{\|I(x, y, t) - A \cdot I(x+u, y+v, t+1)\|^2}_{\text{data term}} + \underbrace{\sigma_s (\|\nabla u\|^2 + \|\nabla v\|^2)}_{\text{spatial reg}} + \underbrace{\sigma_c \|\nabla A\|^2}_{\text{color reg}} dxdy$$





# Some Generative Adversarial Networks

## iGAN





# Some Generative Adversarial Networks

## iGAN



Too good to be true?



# Tricky to train



Original



Generated



The major roadblocks with GANs:  
various optimization issues



# Causes of optimization issues

- Non convergence
- Simultaneous updates require a careful balance between the two players
- There is a stationary point but no guarantee of reaching it
- Mode collapse
- Adversarial optimization is a more general, harder problem than single-player optimization
- Discriminator is highly nonlinear, gradient tends to be noisy or non-informative



# Common failures

## Mode collapse

Generated points tend to "herd" to probable regions, causing "mode collapse".

"Mode collapse" GAN generates a subspace really well but doesn't cover the entire real distribution. For example, train on MNIST and it only generates threes and eights.

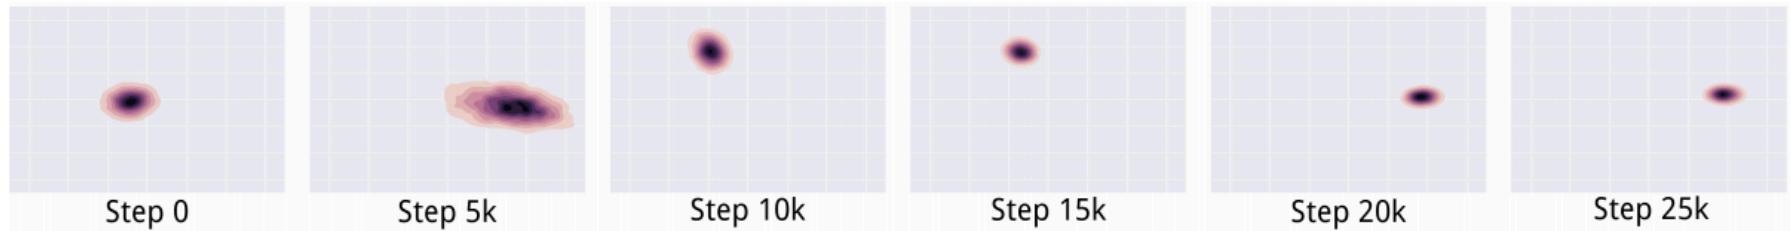
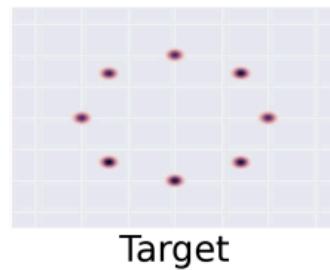
<https://www.youtube.com/watch?v=ktxhiKhWoEE>



# Common failures

Looks like “mode collapse”

Sometimes GANs enter into clear cycles. They seem to generate a single digit relatively well, then start generating a different digit, etc. Looks like “mode collapse” on a rotating set of samples, but it does not differentiate.





# Common failures

## Hard to describe failures

Sometimes hard to describe failures, but videos like this are relatively typical.

<https://www.youtube.com/watch?v=D5akt32hsCQ>

# Optimization Techniques



# Optimization techniques

- There are many tricks to train them better but some of those tricks do away with what makes GANs so special
- Not every trick works all the time or in combination with other tricks
- Most papers claim to have the golden bullet
- Best current solution is really a combination of techniques

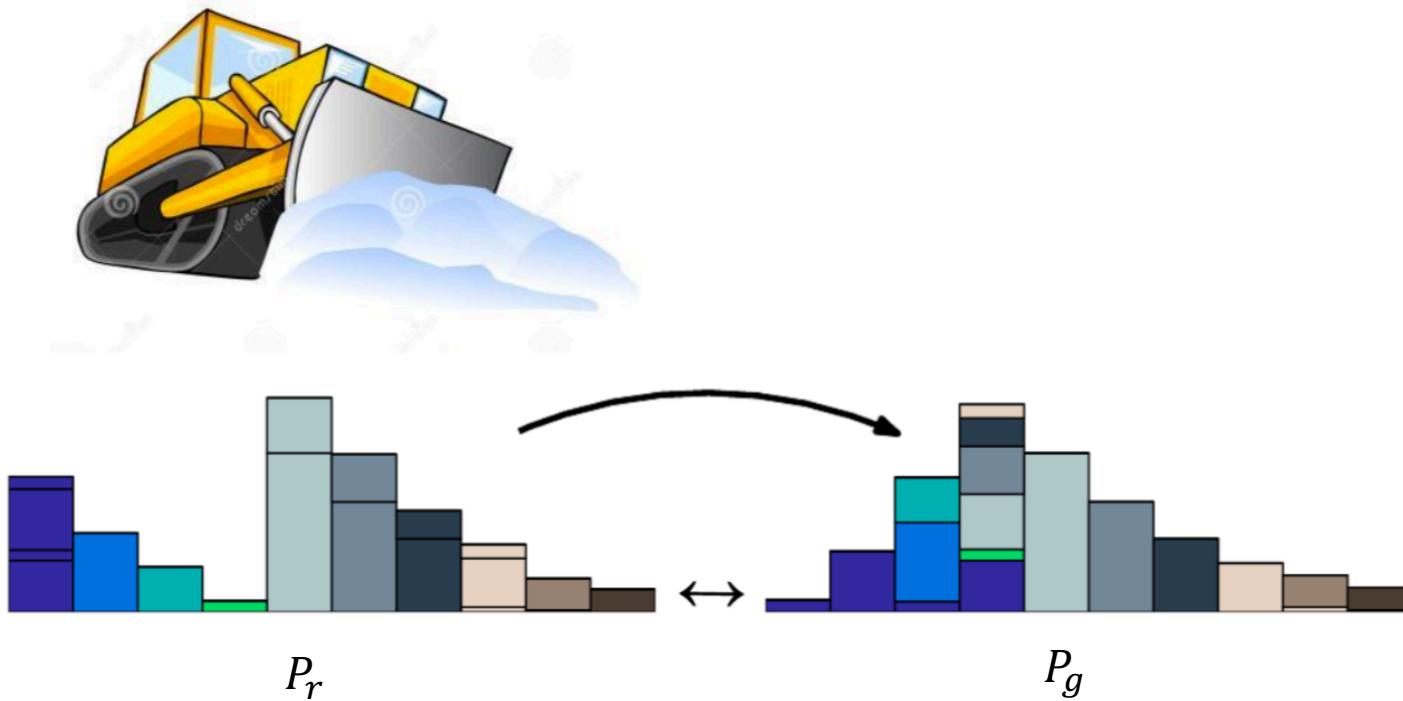


# Optimization techniques

- Unrolled GANs, Metz et al. (2016)
- Improved Techniques for Training GANs, Salimans et al. (2016)
- Least-Squares GAN, Mao et al. (2016)
- Amortised MAP Inference, Kaae Sønderby et al. (2016)
- EBGAN, Zhao et al. (2016)
- WGAN, Arjovsky et al. (2017)
- WGAN-GP, Gulrajani et al. (2017)
- GAN optimization is locally stable, Nagarajan and Kolter (2017)
- Numerics of GANs, Mescheder et al. (2017)
- DRAGAN, Kodali et al. (2017)
- Progressive GAN, Karras et al. (2017)
- Lagrangian GAN, Chen et al. (2018)



# Wasserstein GAN (WGAN)



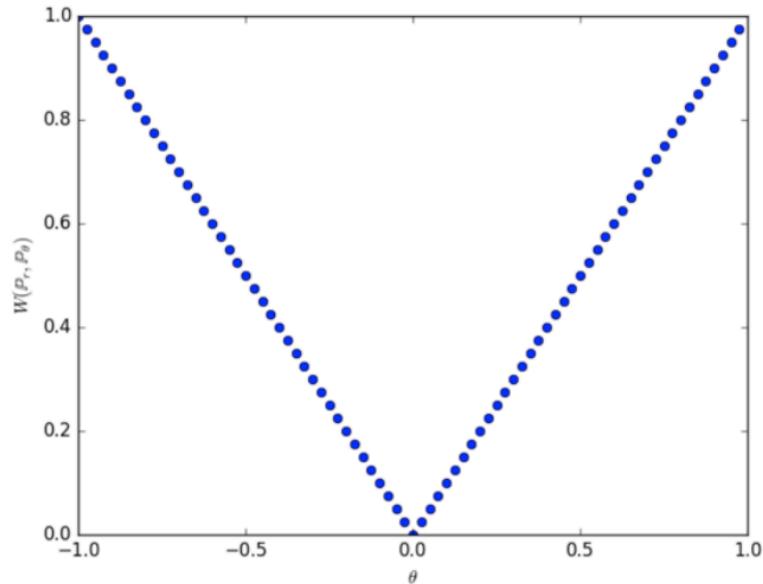
Wasserstein-1 Distance (**Earth-Mover Distance**):

$$W(\mathbb{P}_r, \mathbb{P}_g) = \inf_{\gamma \in \Pi(\mathbb{P}_r, \mathbb{P}_g)} \mathbb{E}_{(x,y) \sim \gamma} [\|x - y\|]$$

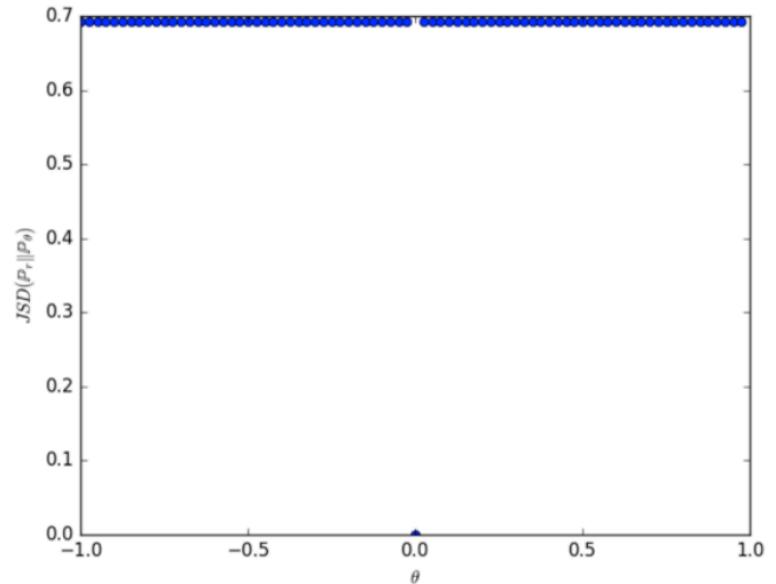


# Wasserstein GAN (WGAN)

EM vs JS



Wasserstein Distance



JS Divergence

$$JS(\mathbb{P}_r, \mathbb{P}_g) = KL(\mathbb{P}_r \parallel \mathbb{P}_m) + KL(\mathbb{P}_g \parallel \mathbb{P}_m)$$



# Wasserstein GAN (WGAN)

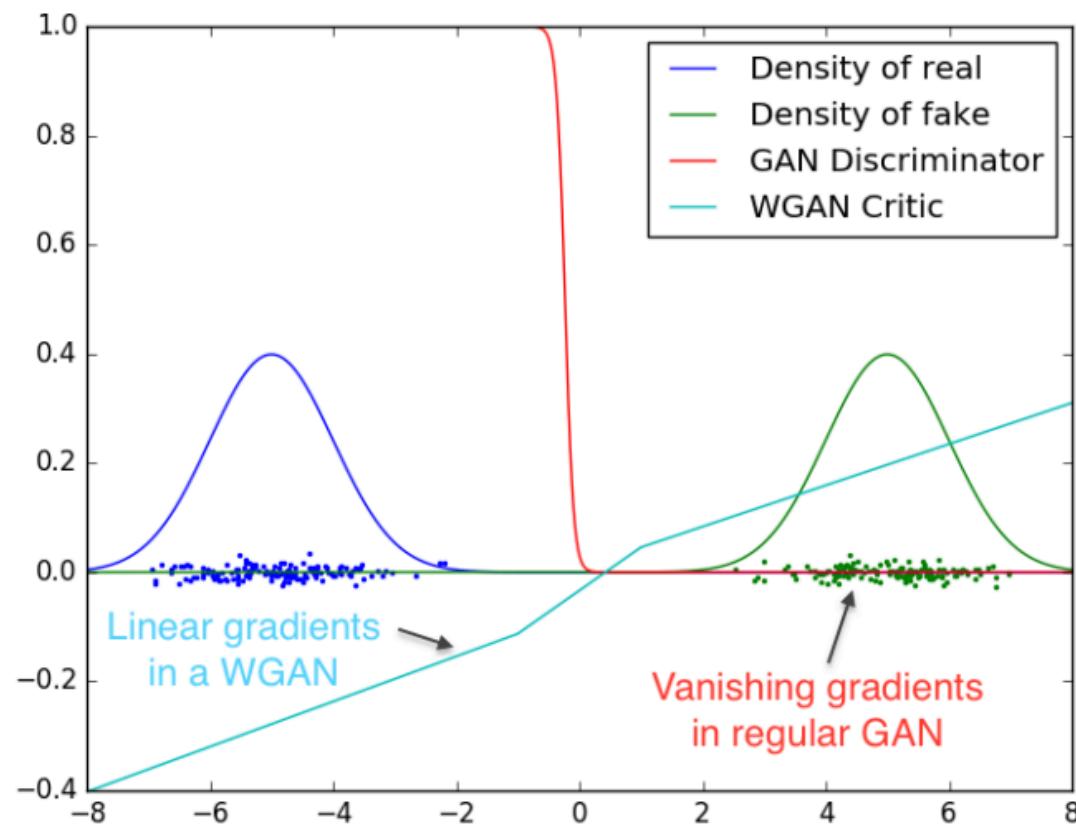


Figure 2: Optimal discriminator and critic when learning to differentiate two Gaussians. As we can see, the discriminator of a minimax GAN saturates and results in vanishing gradients. Our WGAN critic provides very clean gradients on all parts of the space.



# Wasserstein GAN (WGAN)

- Claims more stability than traditional GAN
- Loss value is more meaningful and corresponds to how good samples are
- Less tuning required in terms of generator/discriminator balance
- In practice, doesn't actually work so well (see Improved WGAN and WGAN-GP)



# Wasserstein GAN (WGAN)

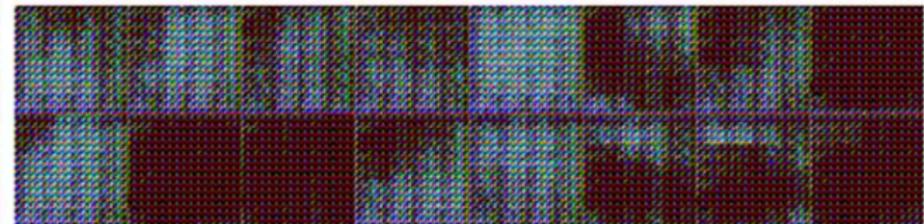
## GAN & DCGAN



WGAN

DCGAN

## Without BN



WGAN

DCGAN



# Progressive GAN

## Architecture

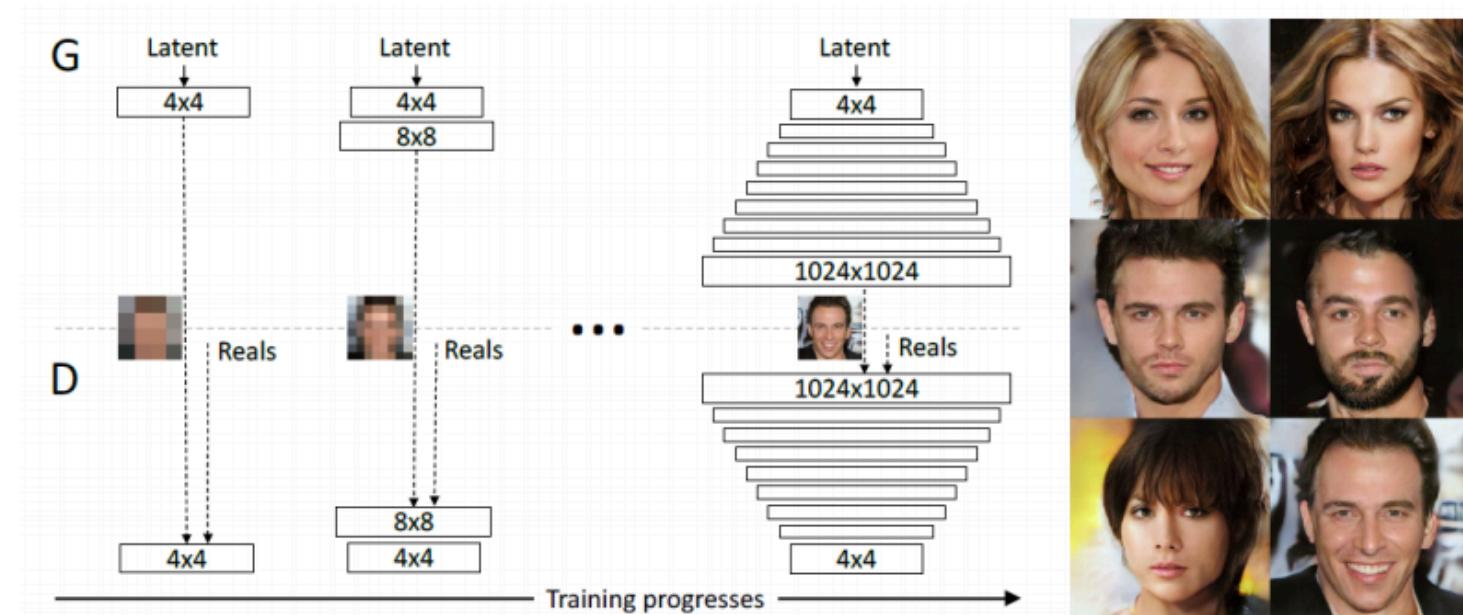


Figure 1: Our training starts with both the generator (G) and discriminator (D) having a low spatial resolution of  $4 \times 4$  pixels. As the training advances, we incrementally add layers to G and D, thus increasing the spatial resolution of the generated images. All existing layers remain trainable throughout the process. Here  $N \times N$  refers to convolutional layers operating on  $N \times N$  spatial resolution. This allows stable synthesis in high resolutions and also speeds up training considerably. On the right we show six example images generated using progressive growing at  $1024 \times 1024$ .



# Progressive GAN

Video demo: <https://www.youtube.com/watch?v=XOxxPcy5Gr4>



Figure 5:  $1024 \times 1024$  images generated using the CELEBA-HQ dataset. See Appendix F for a larger set of results, and the accompanying video for latent space interpolations.

# Image-to-Image Translation



# Image-to-Image Translation



- pix2pix
- CycleGAN
- UNIT
- MUNIT



# pix2pix

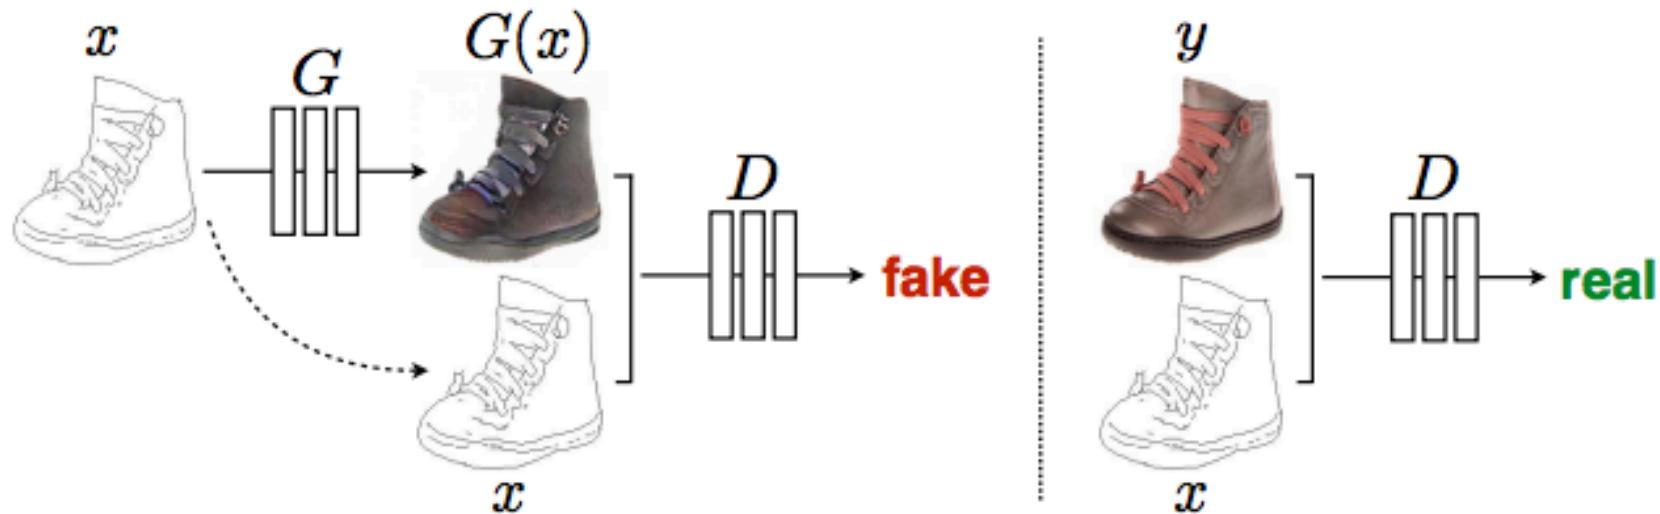


Figure 2: Training a conditional GAN to map edges→photo. The discriminator,  $D$ , learns to classify between fake (synthesized by the generator) and real {edge, photo} tuples. The generator,  $G$ , learns to fool the discriminator. Unlike an unconditional GAN, both the generator and discriminator observe the input edge map.



# pix2pix

## Generator

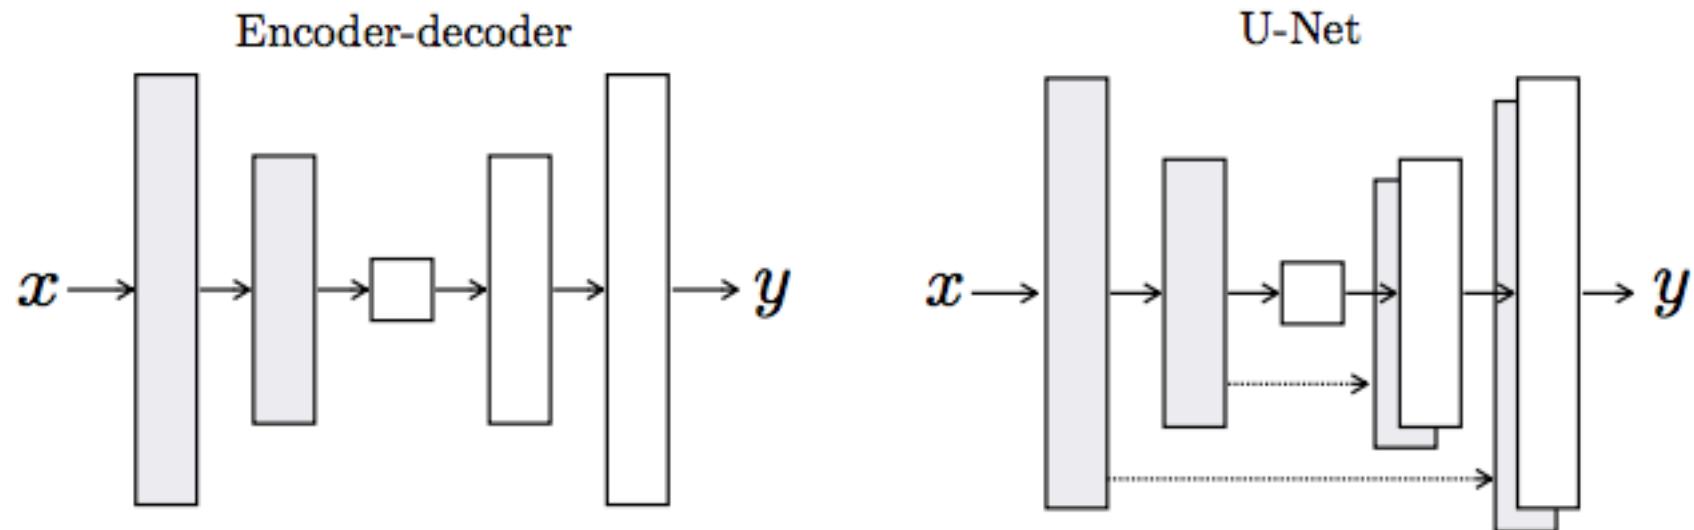


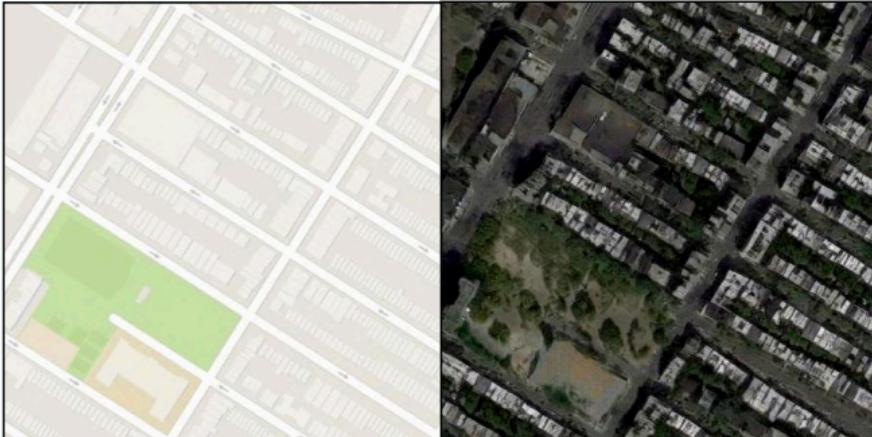
Figure 3: Two choices for the architecture of the generator. The “U-Net” [49] is an encoder-decoder with skip connections between mirrored layers in the encoder and decoder stacks.



# pix2pix

Map  $\leftrightarrow$  Aerial photo

Map to aerial photo



Aerial photo to map



input

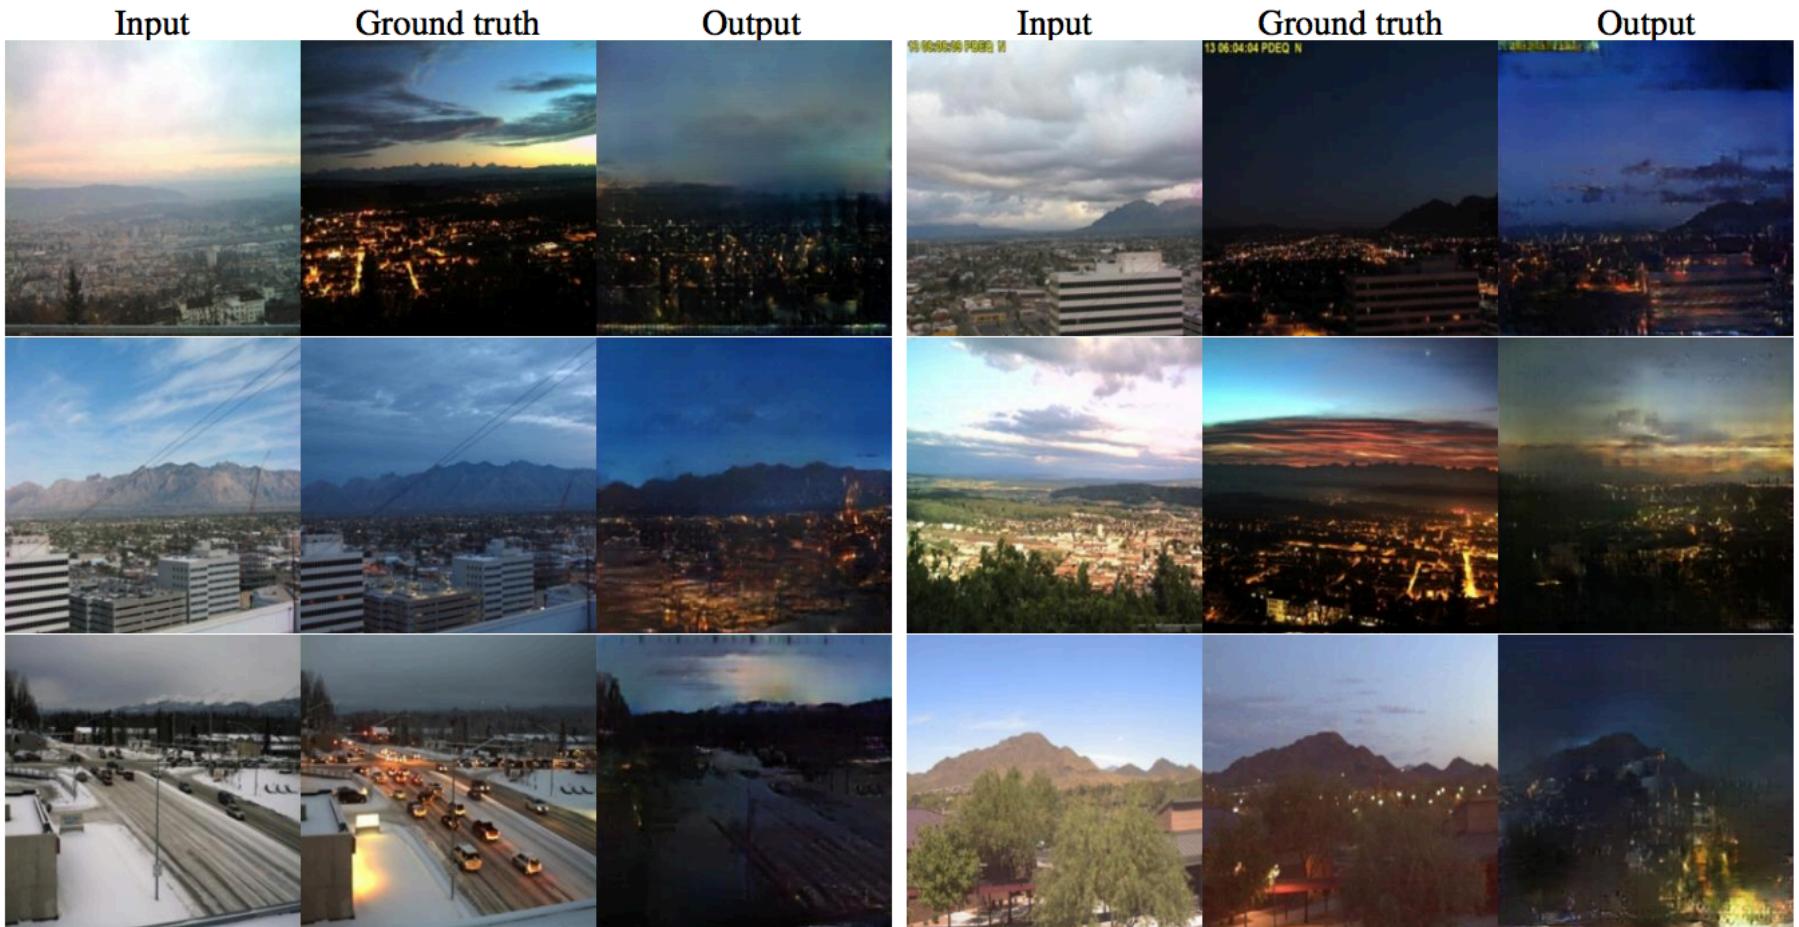
output

output



# pix2pix

Day  $\leftrightarrow$  Night

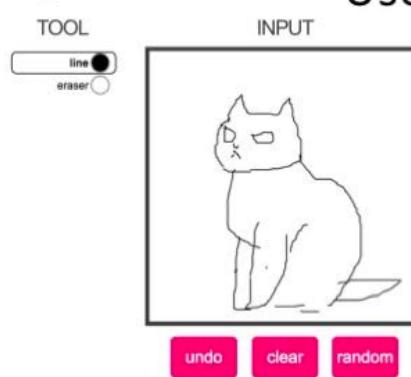




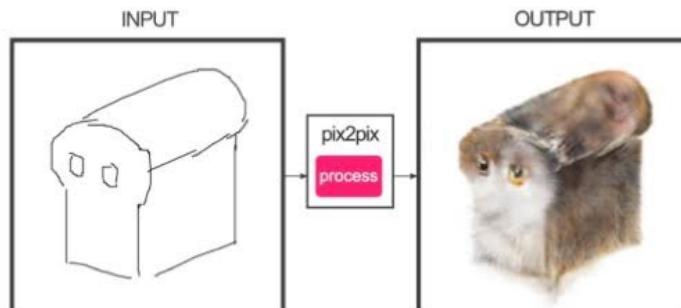
# pix2pix

Online demo: <https://affinelayer.com/pixsrv/>

edges2cats



@gods\_tail



Results



Vitaly Vidmirov @vvivid



@ka92



# CycleGAN

Paired

$$x_i \quad y_i$$



Unpaired

$X$

$Y$





# CycleGAN

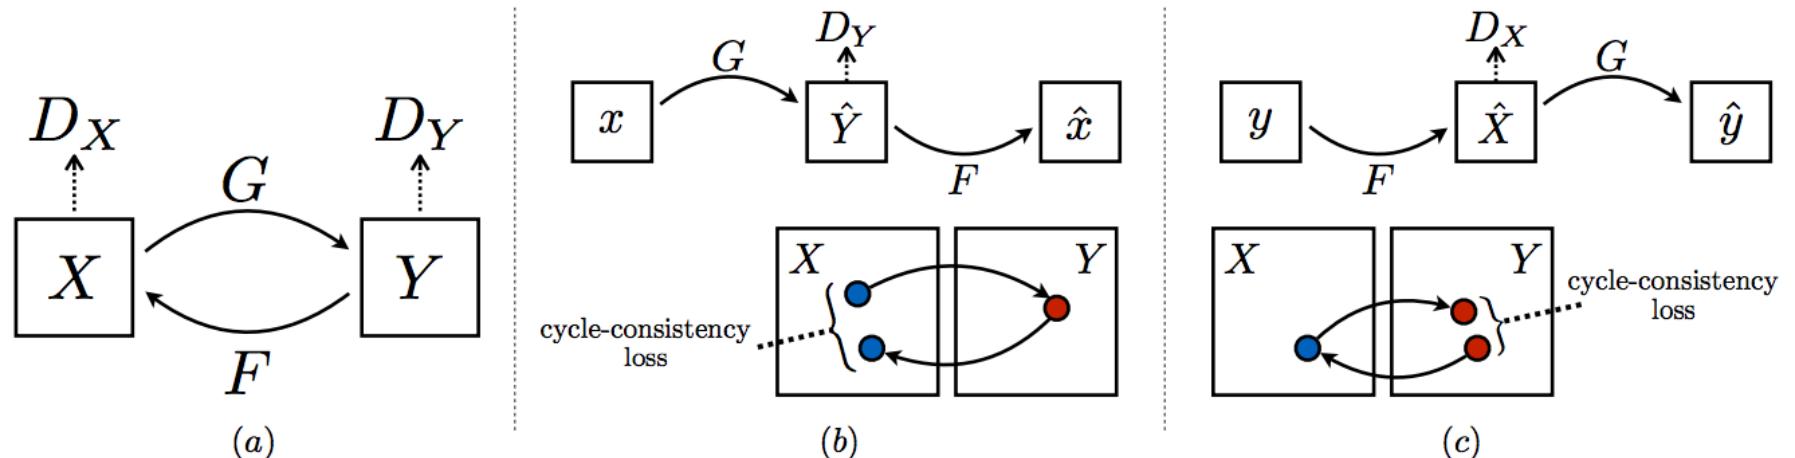


Figure 3: (a) Our model contains two mapping functions  $G : X \rightarrow Y$  and  $F : Y \rightarrow X$ , and associated adversarial discriminators  $D_Y$  and  $D_X$ .  $D_Y$  encourages  $G$  to translate  $X$  into outputs indistinguishable from domain  $Y$ , and vice versa for  $D_X$  and  $F$ . To further regularize the mappings, we introduce two *cycle consistency losses* that capture the intuition that if we translate from one domain to the other and back again we should arrive at where we started: (b) forward cycle-consistency loss:  $x \rightarrow G(x) \rightarrow F(G(x)) \approx x$ , and (c) backward cycle-consistency loss:  $y \rightarrow F(y) \rightarrow G(F(y)) \approx y$

$$\begin{aligned} \mathcal{L}_{\text{cyc}}(G, F) = & \mathbb{E}_{x \sim p_{\text{data}}(x)} [\|F(G(x)) - x\|_1] \\ & + \mathbb{E}_{y \sim p_{\text{data}}(y)} [\|G(F(y)) - y\|_1]. \end{aligned}$$



# CycleGAN

horse  $\leftrightarrow$  zebra





# CycleGAN

photo  $\leftrightarrow$  drawing





# CycleGAN

Not perfect

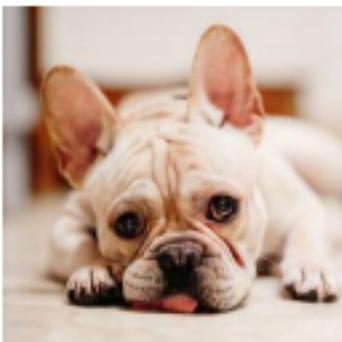
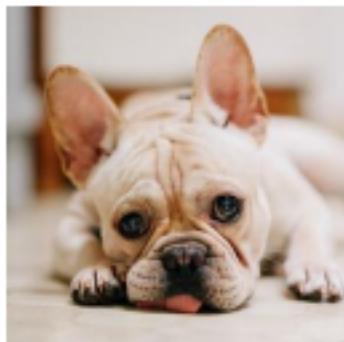
Input



Output



apple → orange

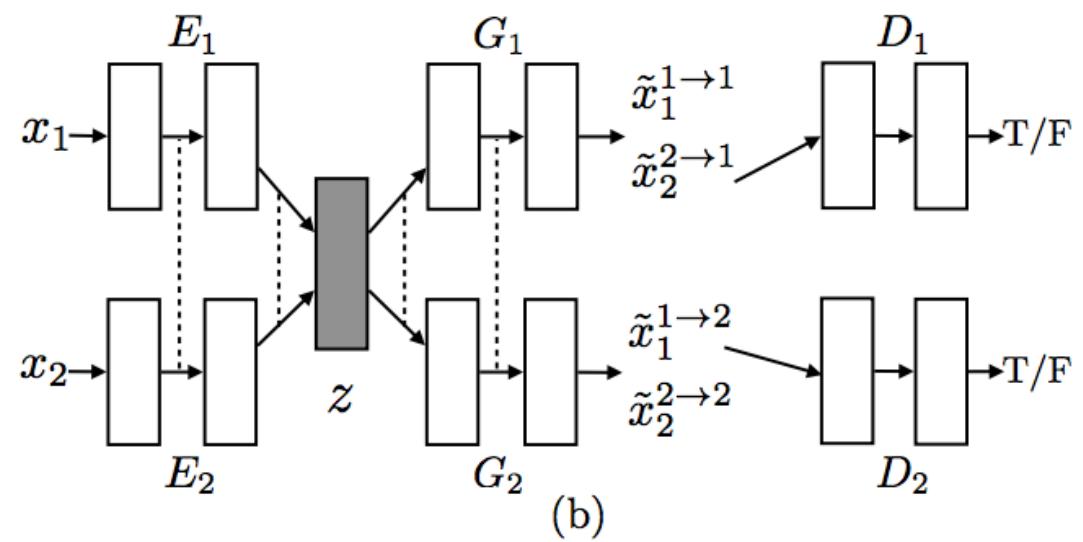
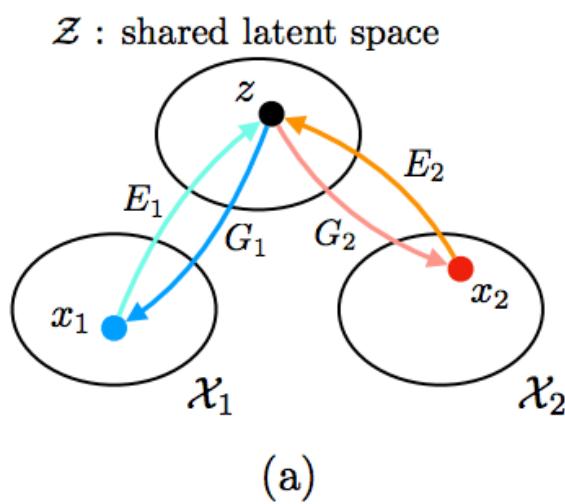


dog → cat



Horse → Zebra

# UNIT





# UNIT

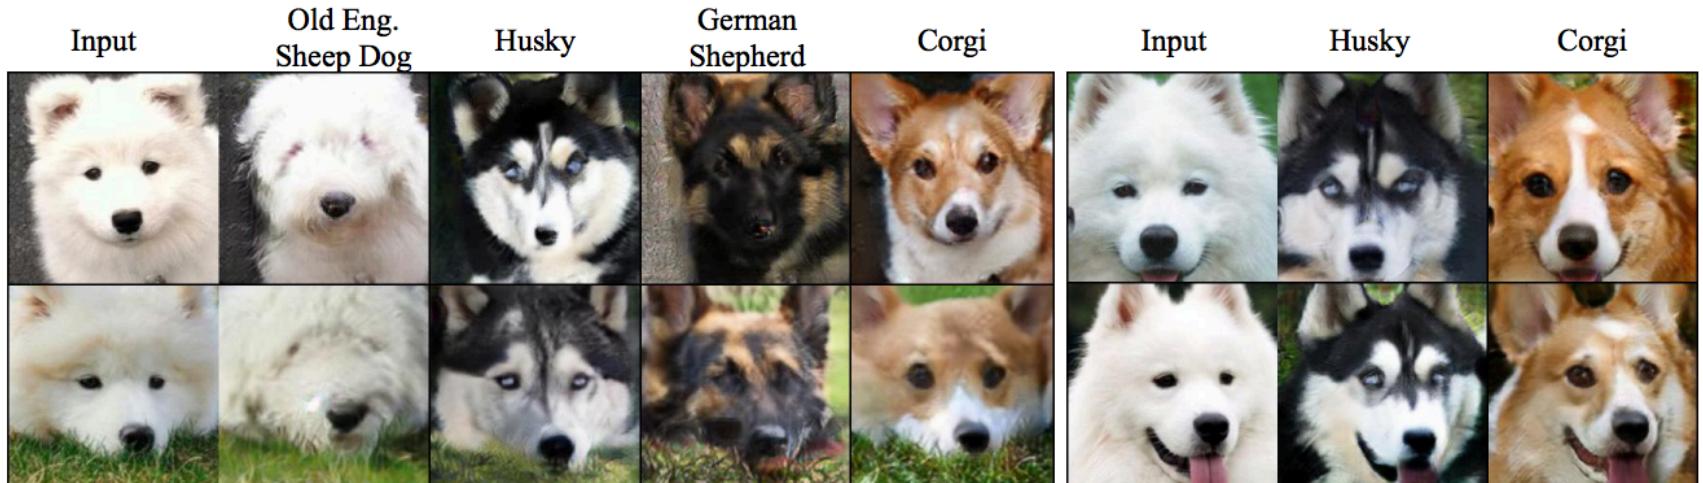


Figure 4: Dog breed translation results.

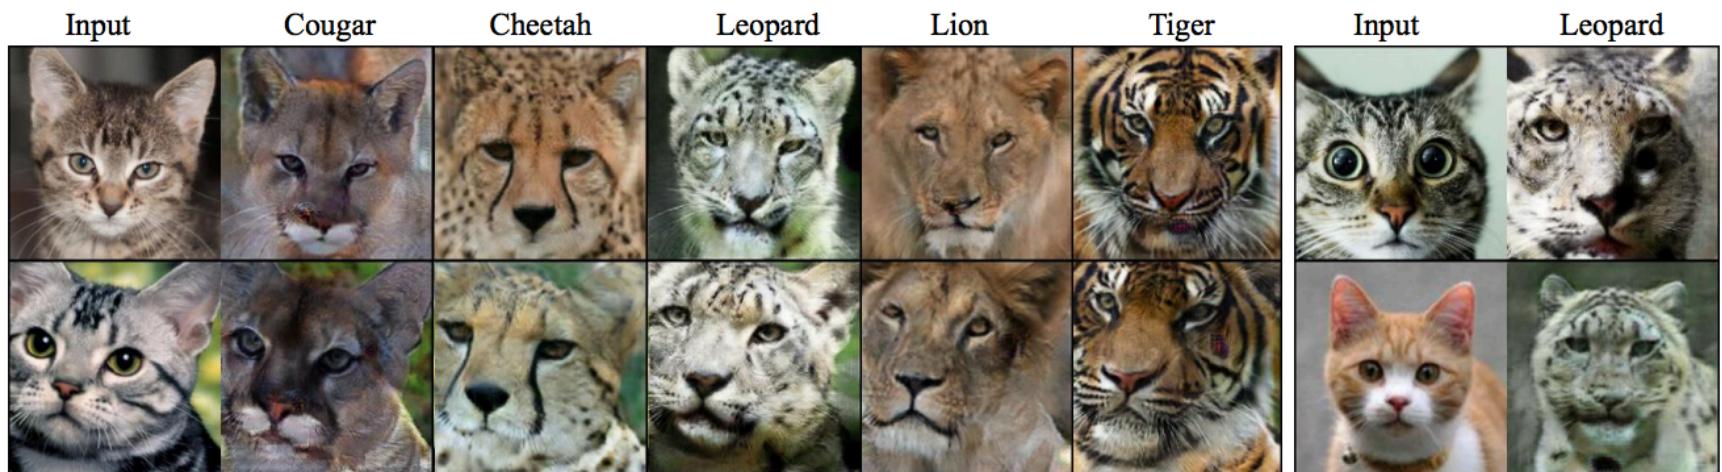
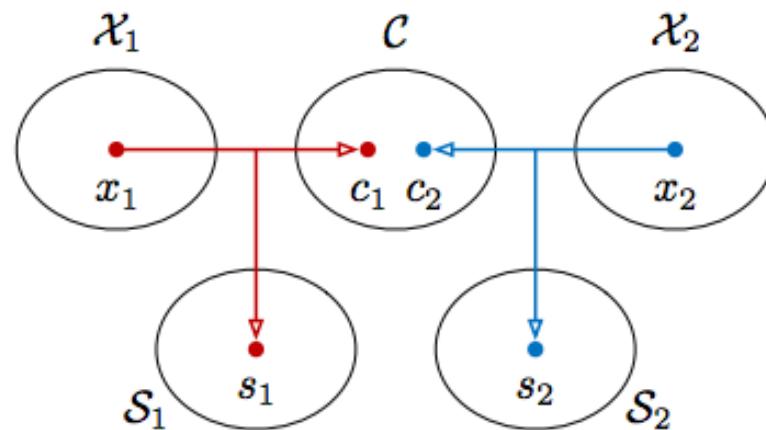


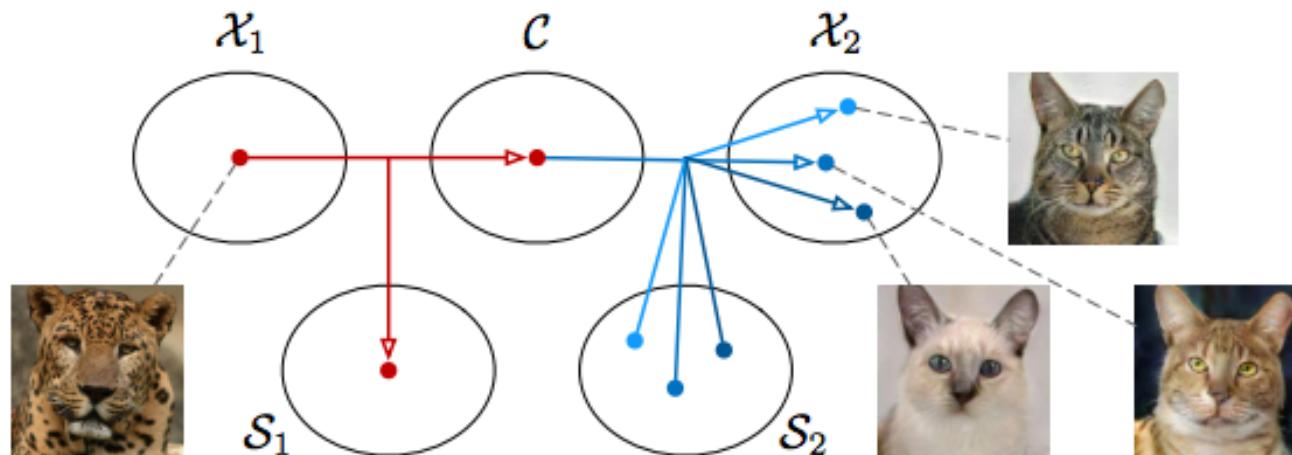
Figure 5: Cat species translation results.



# MUNIT



(a) Auto-encoding



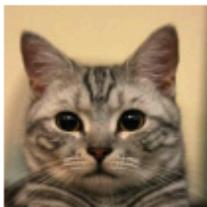
(b) Translation



# MUNIT

Video demo: <https://youtu.be/ab64TWzWn40>

Input



Sample translations

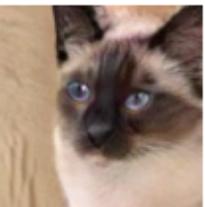


(a) house cats → big cats

Input



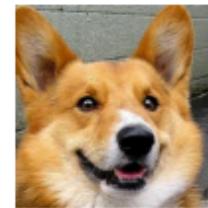
Sample translations



(b) big cats → house cats



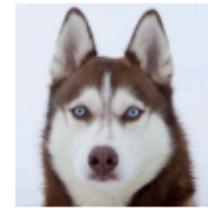
(c) house cats → dogs



(d) dogs → house cats



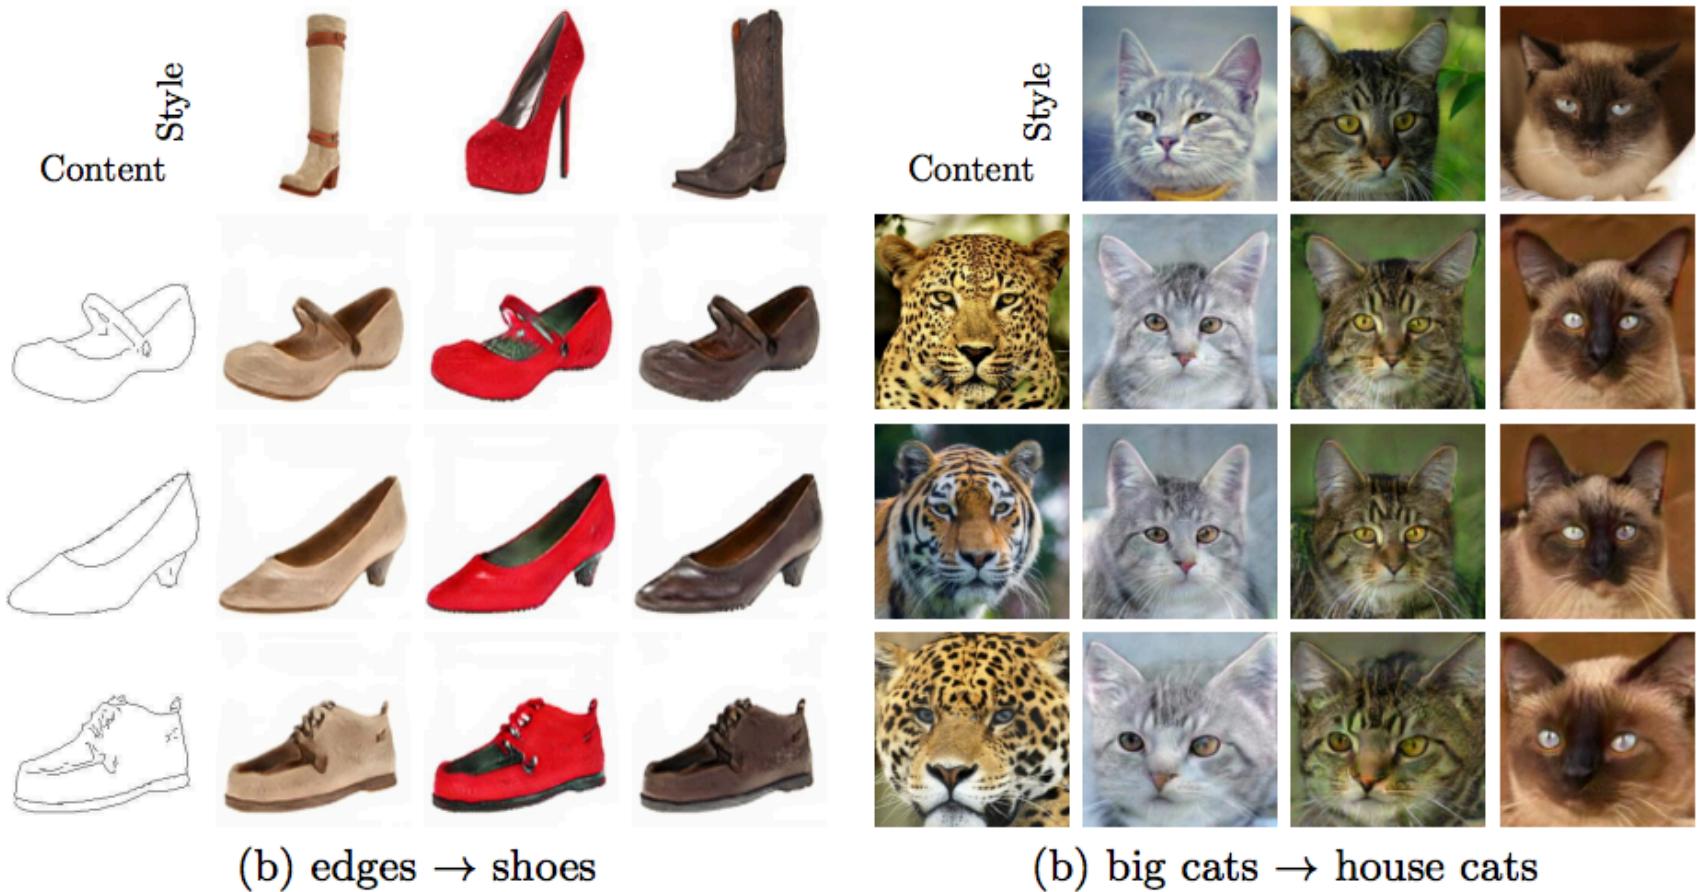
(e) big cats → dogs



(f) dogs → big cats



# MUNIT



**Fig. 9.** Example-guided image translation. Each row has the same content while each column has the same style.

Tons of other applications



# Tons of applications

- Face Synthesis
- Anime Illustration Generation
- Colorize Pictures
- Photo/Video Editing
- Edge to Drawing
- Text to Image
- Image Inpainting
- Image Restoration
- Super Resolution
- Image Deraining/Desnowing
- 2D to 3D
- and so on ...



# Style Transfer

字 種 成 東 字 推  
符 利 對 亞 型 斷  
到 用 抗 語 進 的  
字 條 網 言 行 新  
符 件 絡 字 自 方  
一 生 對 體 動 法





# Chinese Poem Generation

- 曾中书惟静，汝道卖晚寒。池好可河夫，高怀礼呼山。
- 风下带红云，目烟坐井缘，头老起香寥，谢斜影日生。
- 一机不识飞，各有何处人。试知人无处，不复行相人。
- 清暑浮云晓，天水征群近。流流只是入，年事比玩庭。



# Face aging





# Cross-generation Face Recognition





# Finding Abducted Children



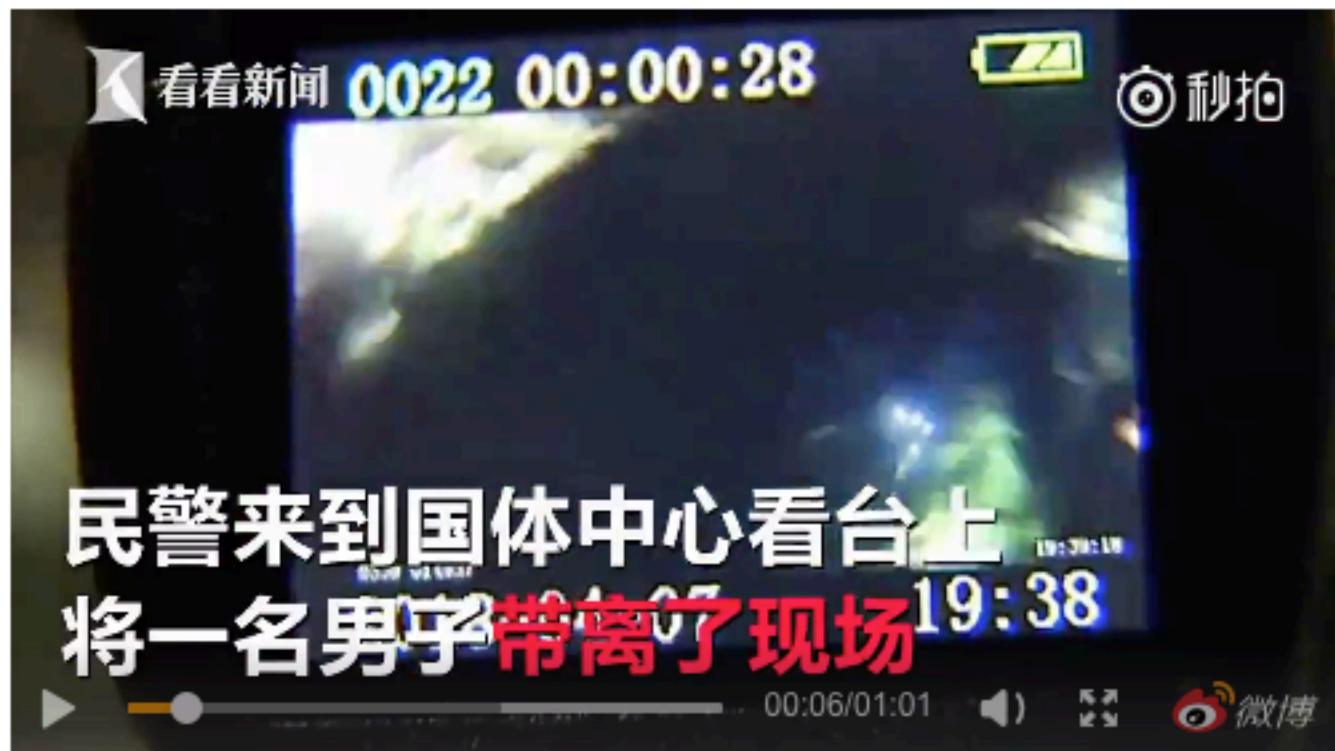
5岁时                    8岁时

XX , 1984年在重庆出生 , 为家中长子长孙 , 1990年在重庆石柱县丢失 , 后被拐到福建泉州 , 丢失27年 , 现33岁。



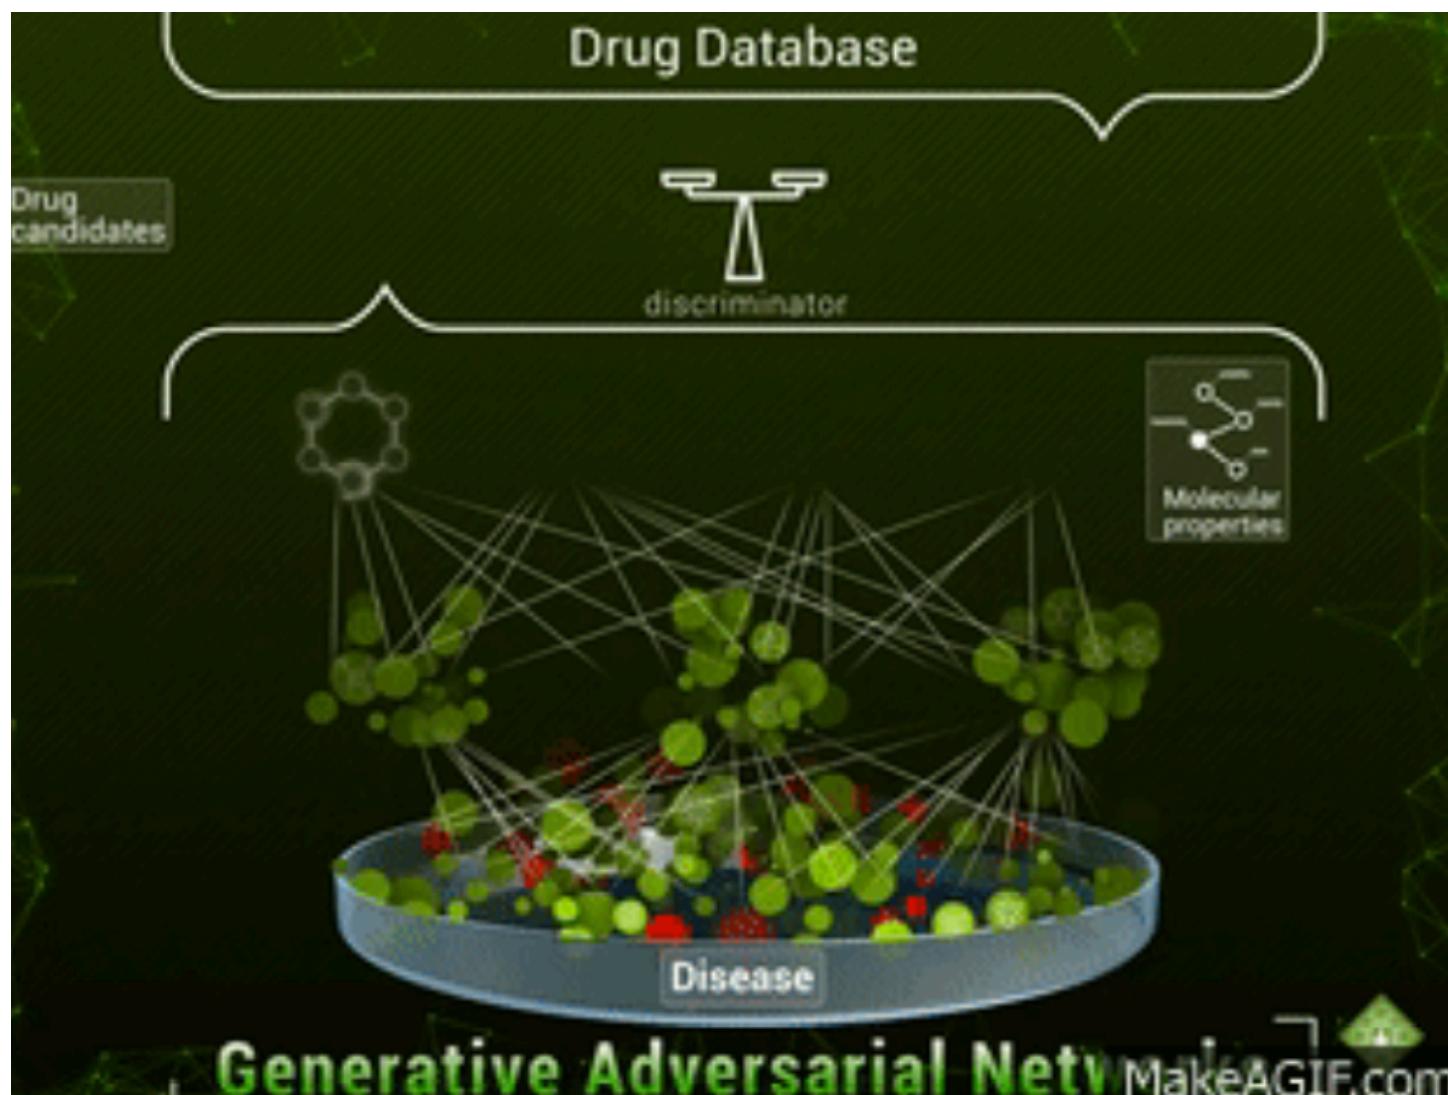
# Image Restoration & Face Recognition (maybe GANs are not used)

【意不意外？逃犯听张学友的演唱会 被人像识别认出落网】7日晚，[张学友南昌演唱会](#)开演后没多久，民警就在演出中心的看台上将一名男子带离现场。据警方介绍，男子姓敖，31岁，江西人，被列为网上追逃。前往现场看演唱会时被智慧安保人像识别功能锁定，让民警在茫茫人海中找到他。[看看新闻Kn...](#)





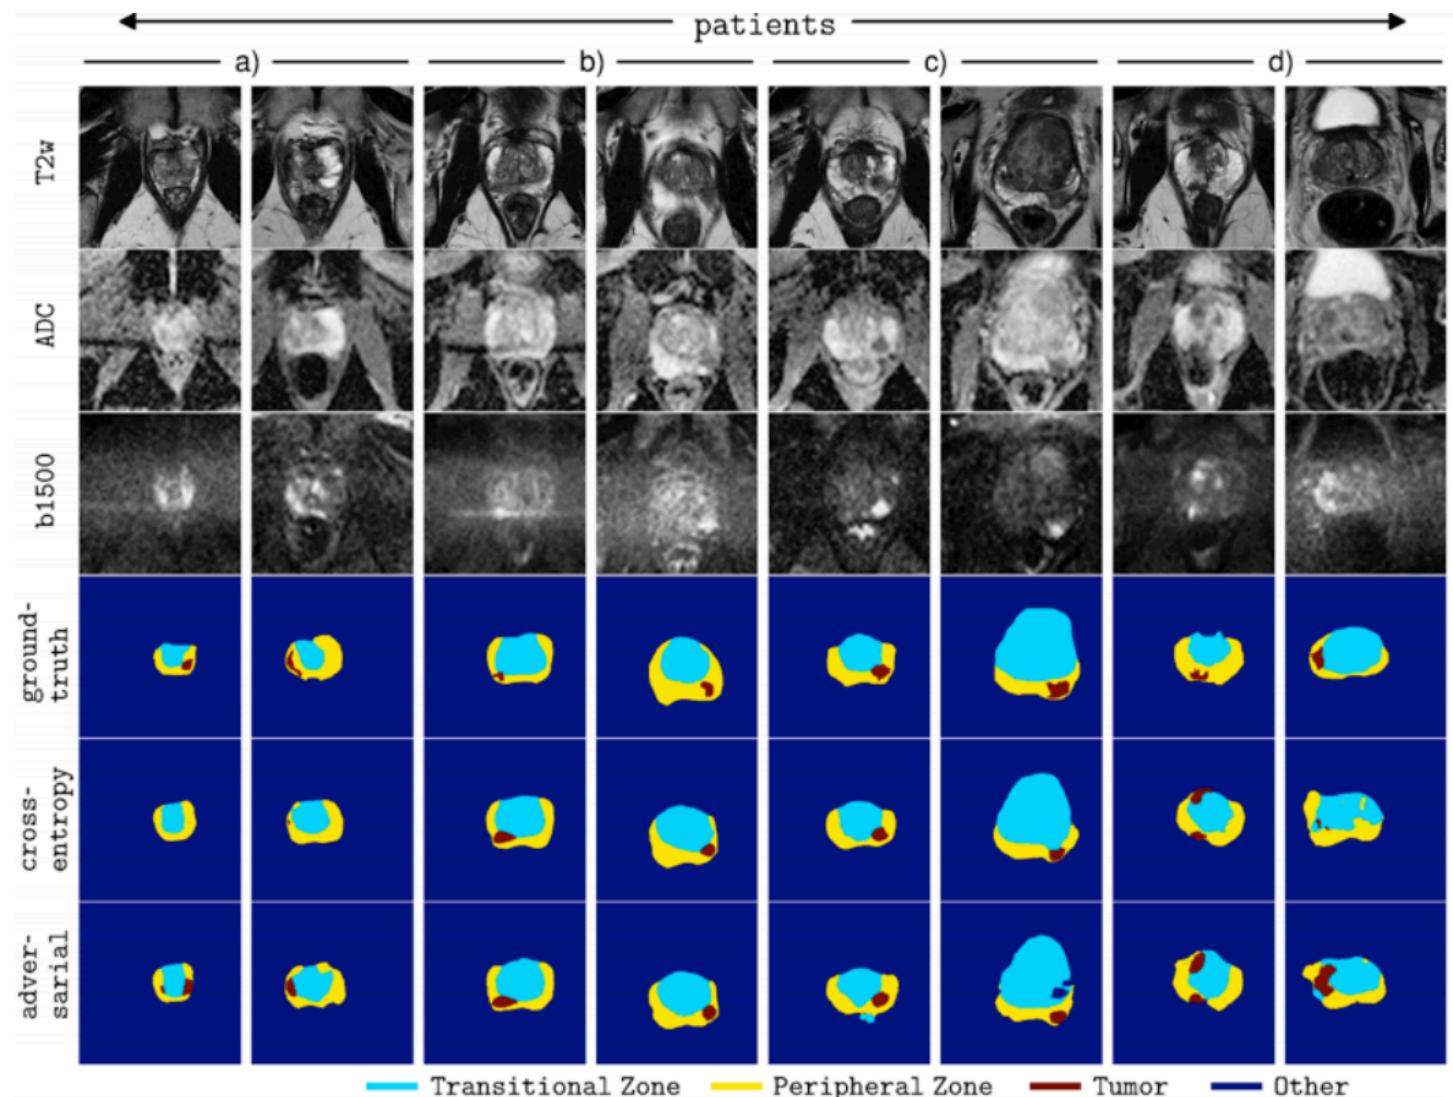
# Drug Discovery



Kadurin *et al.*, "The cornucopia of meaningful leads: Applying deep adversarial autoencoders for new molecule development in oncology". Oncotarget. 2017.



# Automatic Cancer Detection





# Take away message

GAN have a lot of potential

- Flexible to support a variety of problems
- Potential to represent meaningful metrics

GAN optimization, architecture, etc. is still an active area of research

# TensorFlow Tutorial



# What is TensorFlow?

“Open source software library for numerical computation using data flow graphs”



# Why TensorFlow?



**Denny Britz** @dennybritz · 25 Dec 2017



I'm going through my newsletters to write up a year-end summary of developments and achievements in AI.

Fun fact: Almost every week, a company released a new generic or task-specific Deep Learning “framework” 😅

- TensorFlow
- Keras
- PyTorch
- MXNet
- Caffe
- Torch
- CNTK
- PaddlePaddle
- .....



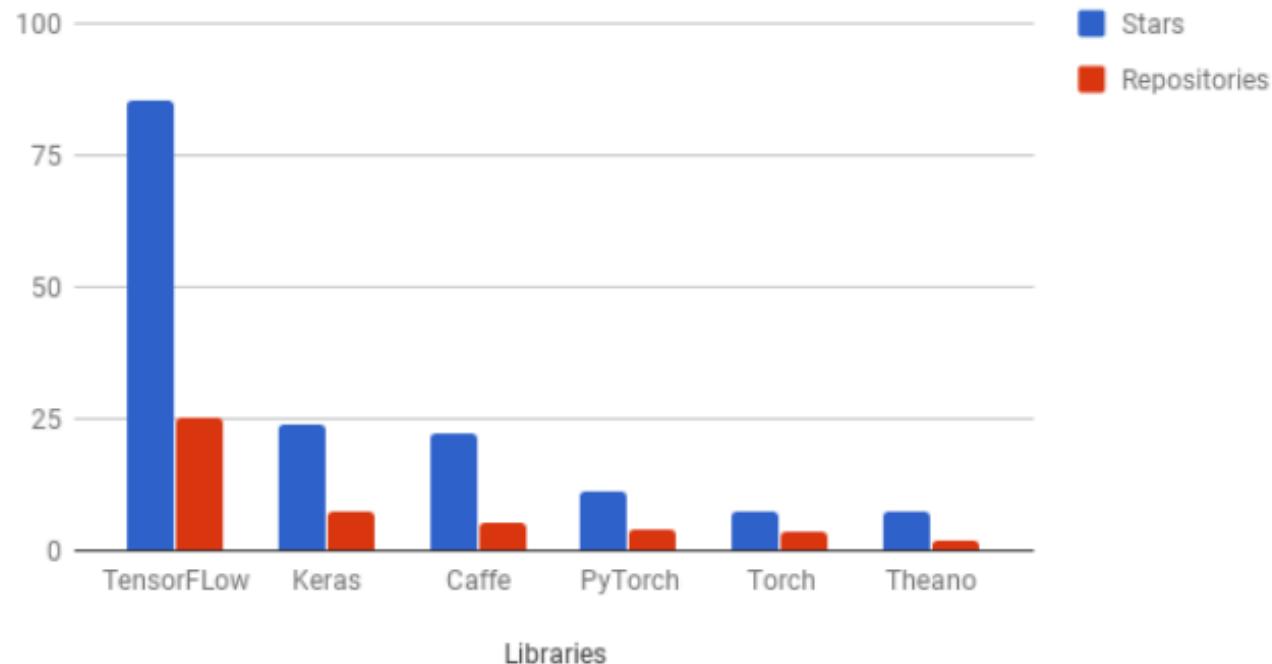
# Why TensorFlow?

Flexibility + Scalability

Popularity



Stars and Repositories





# Companies using TensorFlow



kakao



ZTE



3DR



ARM



JD.COM 京东





# Getting started

```
import tensorflow as tf
```



# Graphs and Sessions

## Data Flow Graphs

TensorFlow separates definition of computations from their execution

Phase 1: assemble a graph

Phase 2: use a session to execute operations in the graph.



# Graphs and Sessions

## Data Flow Graphs

TensorFlow separates definition of computations from their execution

Phase 1: assemble a graph

Phase 2: use a session to execute operations in the graph.

Changed in **eager mode!!!**



# What's a tensor?

**An n-dimensional array**

- 0-d tensor: scalar  
(number)
- 1-d tensor: vector
- 2-d tensor: matrix
- and so on



# Data Flow Graphs

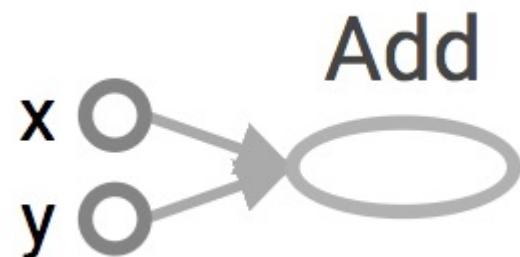
```
import tensorflow as tf  
  
a = tf.add(3, 5)
```



# Data Flow Graphs

```
import tensorflow as tf  
a = tf.add(3, 5)
```

Visualize with **tensorboard**





# Data Flow Graphs

```
import tensorflow as tf  
a = tf.add(3, 5)
```

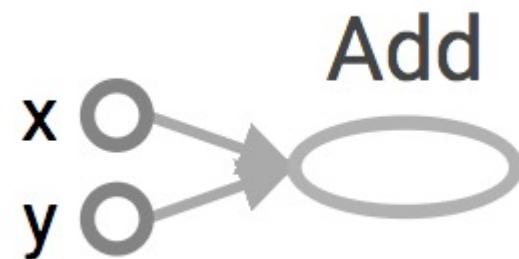
Why x, y?

TF automatically names the nodes when you don't explicitly name them.

```
x = 3
```

```
y = 5
```

Visualize with **tensorboard**

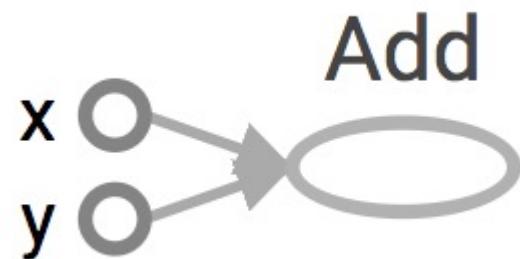




# Data Flow Graphs

```
import tensorflow as tf  
a = tf.add(3, 5)
```

Visualize with **tensorboard**



Nodes: operators, variables, and constants

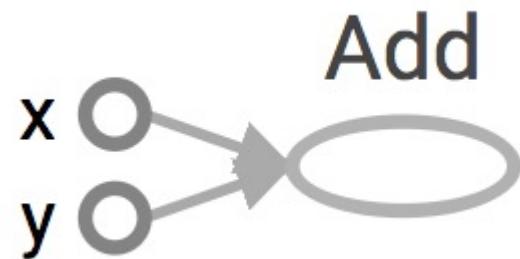
Edges: tensors



# Data Flow Graphs

```
import tensorflow as tf  
a = tf.add(3, 5)  
print(a)
```

Visualize with **tensorboard**



```
>> Tensor("Add:0", shape=(), dtype=int32)  
(Not 8)
```



# How to get the value of $a$ ?

Create a **session**, assign it to variable sess so we can call it later

Within the session, evaluate the graph to fetch the value of  $a$



# How to get the value of *a*?

Create a **session**, assign it to variable sess so we can call it later

Within the session, evaluate the graph to fetch the value of *a*

```
import tensorflow as tf  
a = tf.add(3, 5)  
sess = tf.Session()  
print(sess.run(a))  
sess.close()
```

>> 8





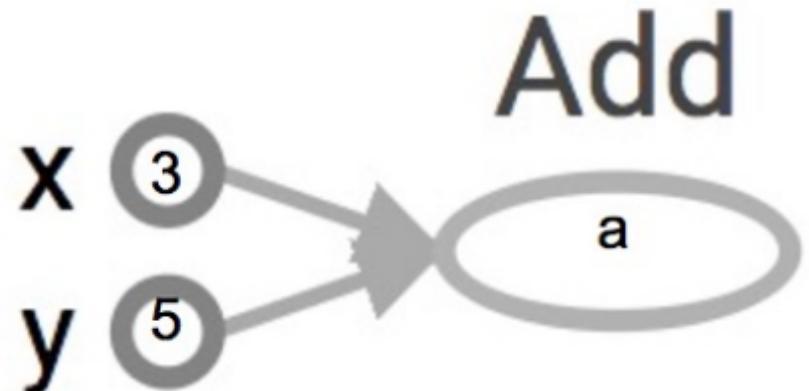
# How to get the value of *a*?

Create a **session**, assign it to variable sess so we can call it later

Within the session, evaluate the graph to fetch the value of *a*

```
import tensorflow as tf  
a = tf.add(3, 5)  
with tf.Session() as sess:  
    print(sess.run(a))
```

>> 8





## *tf.Session()*

A Session object encapsulates the environment in which Operation objects are executed, and Tensor objects are evaluated.

Session will also allocate memory to store the current values of variables.

Thank you!